

CS599 (Deep Learning)

Homework – 3

1. Python Code:

```
import pandas as pd
from sklearn.model_selection import KFold, GridSearchCV

import matplotlib
matplotlib.use("agg")

#preparing zip data for binary classification
zip_df = pd.read_csv("zip.test.gz", sep = " ", header = None)
zip_label_col_num = 0
zip_label_vec = zip_df.iloc[:, zip_label_col_num]
is_01 = zip_label_vec.isin([0,1])
zip_01_df = zip_df.loc[is_01, :]
is_label_col = zip_01_df.columns == zip_label_col_num
zip_features = zip_01_df.iloc[:, ~is_label_col]
zip_labels = zip_01_df.iloc[:, is_label_col]

#preparing spam data for binary classification
spam_df = pd.read_csv("spam.data", sep= " ", header = None)
spam_label_col_num = -1
spam_label_vec = spam_df.iloc[:, spam_label_col_num]
spam_is_01 = spam_label_vec.isin([0,1])
spam_01_df = spam_df.loc[spam_is_01, :]
spam_features = spam_df.iloc[:, :spam_label_col_num]
spam_labels = spam_df.iloc[:, spam_label_col_num]

import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

data_dict = {
    "zip" : (zip_features, zip_labels),
    "spam" : (spam_features, spam_labels)
}

accuracy_data_frames = []
for data_name, (data_features, data_labels) in data_dict.items():
    kf = KFold(n_splits = 3, shuffle = True, random_state = 3)
    enum_obj = enumerate(kf.split(data_features))
    for fold_num, (train_index, test_index) in enum_obj:
```

```

X_train, X_test = np.array(data_features.iloc[train_index]), np.array(data_features.iloc[test_index])

y_train, y_test = np.ravel(data_labels.iloc[train_index]), np.ravel(data_labels.iloc[test_index])


#K-nearest neighbors
knn = KNeighborsClassifier()
hp_parameters = {"n_neighbors": list(range(1,21))}
grid = GridSearchCV(knn, hp_parameters, cv = 5)
grid.fit(X_train, y_train)
best_n_neighbors = grid.best_params_['n_neighbors']
print("Best N-Neighbors = ", best_n_neighbors)
knn = KNeighborsClassifier(n_neighbors = best_n_neighbors)
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)


#Logistic Regression
pipe = make_pipeline(StandardScaler(), LogisticRegressionCV(cv=5, max_iter=2000))
pipe.fit(X_train, y_train)
lr_pred = pipe.predict(X_test)
y_train_series = pd.Series(y_train)
most_frequent_class = y_train_series.value_counts().idxmax()
print("Most Frequent Class = ", most_frequent_class)


#create a featureless baseline
featureless_pred = np.repeat(most_frequent_class, len(y_test))


#store predict data in dict
pred_dict = {'nearest neighbors': knn_pred,
             'linear_model': lr_pred,
             'featureless': featureless_pred}
test_accuracy = {}


for algorithm, predictions in pred_dict.items():
    accuracy = accuracy_score(y_test, predictions)
    test_accuracy[algorithm] = accuracy


for algorithm, accuracy in test_accuracy.items():
    print(f"{algorithm} Test Accuracy: {accuracy * 100}")
    accuracy_df = pd.DataFrame({
        "data_set": [data_name],
        "fold_id": [fold_num],
        "algorithm": [algorithm],
        "accuracy": [test_accuracy[algorithm]]})
    accuracy_data_frames.append(accuracy_df)
    print(f"*****End of
{data_name}({fold_num})*****")
total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
print(total_accuracy_df)

```

```
import plotnine as p9

gg = p9.ggplot(total_accuracy_df, p9.aes(x='accuracy', y='algorithm', fill='data_set'))+\
    p9.facet_grid('~data_set') + p9.geom_point()

gg.save("Output.png")
```

2. Output:

```
>>> for data_name, (data_features, data_labels) in data_dict.items():
...     kf = KFold(n_splits = 3, shuffle = True, random_state = 3)
...     ...
...     print(f"*****End of
{data_name}{{fold_num}}*****")
```

```
Best N-Neighbors = 1
Most Frequent Class = 0
nearest neighbors Test Accuracy: 100.0
linear_model Test Accuracy: 99.51923076923077
featureless Test Accuracy: 58.65384615384615
*****End of zip(0)*****
```

```
Best N-Neighbors = 1
Most Frequent Class = 0
nearest neighbors Test Accuracy: 99.51923076923077
linear_model Test Accuracy: 99.03846153846155
featureless Test Accuracy: 57.21153846153846
*****End of zip(1)*****
```

```
Best N-Neighbors = 3
Most Frequent Class = 0
nearest neighbors Test Accuracy: 99.03381642512076
linear_model Test Accuracy: 99.03381642512076
featureless Test Accuracy: 57.00483091787439
*****End of zip(2)*****
```

```
Best N-Neighbors = 3
Most Frequent Class = 0
nearest neighbors Test Accuracy: 79.85658409387223
linear_model Test Accuracy: 91.39504563233378
featureless Test Accuracy: 60.88657105606258
*****End of spam(0)*****
```

```
Best N-Neighbors = 5
Most Frequent Class = 0
nearest neighbors Test Accuracy: 77.90091264667535
linear_model Test Accuracy: 92.63363754889178
featureless Test Accuracy: 60.104302477183836
*****End of spam(1)*****
```

```
Best N-Neighbors = 1
Most Frequent Class = 0
```

nearest neighbors Test Accuracy: 81.99608610567515

linear_model Test Accuracy: 92.8897586431833

featureless Test Accuracy: 60.79582517938682

*****End of spam(2)*****


```
>>> total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
```

```
>>> print(total_accuracy_df)
```

	data_set	fold_id	algorithm	accuracy
0	zip	0	nearest neighbors	1.000000
1	zip	0	linear_model	0.995192
2	zip	0	featureless	0.586538
3	zip	1	nearest neighbors	0.995192
4	zip	1	linear_model	0.990385
5	zip	1	featureless	0.572115
6	zip	2	nearest neighbors	0.990338
7	zip	2	linear_model	0.990338
8	zip	2	featureless	0.570048
9	spam	0	nearest neighbors	0.798566
10	spam	0	linear_model	0.913950
11	spam	0	featureless	0.608866
12	spam	1	nearest neighbors	0.779009
13	spam	1	linear_model	0.926336
14	spam	1	featureless	0.601043
15	spam	2	nearest neighbors	0.819961
16	spam	2	linear_model	0.928898
17	spam	2	featureless	0.607958

```
>>> gg = p9.ggplot(total_accuracy_df, p9.aes(x='accuracy', y='algorithm', fill='data_set'))+\
```

```
... p9.facet_grid('~data_set') + p9.geom_point()
```



3. Summary:

- First, we need to prepare the data such that it contains 0's and 1's in any of the labels, so that we can perform binary classification.
- To do that, we need to remove all non-01 labels from the both datasets.
- Need to create a data dictionary and run a loop over it.
- Use sklearn package to perform KFold validation, GridSearch, KNeighborsClassifier, LogisticRegression.
- Create a prediction dictionary and print all the 3 prediction accuracy. (Nearest Neighbors, Linear Model & Featureless)
- Make a ggplot using geom_point().