

CS599 (Deep Learning)

Homework – 07

1. Python Code:

```
import torch
import pandas as pd
import matplotlib
matplotlib.use("agg")
import numpy as np
import math
import plotnine as p9

from sklearn.model_selection import KFold, GridSearchCV, ParameterGrid
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from collections import Counter

data_set_dict = {"zip": ("zip.test.gz", 0),
                 "spam": ("spam.data", 57)}
data_dict = {}

for data_name, (file_name, label_col_num) in data_set_dict.items():
    data_df = pd.read_csv(file_name, sep=" ", header=None)
    data_label_vec = data_df.iloc[:, label_col_num]
    is_01 = data_label_vec.isin([0, 1])
    data_01_df = data_df.loc[is_01, :]
    is_label_col = data_df.columns == label_col_num
    data_features = data_01_df.iloc[:, ~is_label_col]
    data_labels = data_01_df.iloc[:, is_label_col]
    data_dict[data_name] = (data_features, data_labels)

spam_features, spam_labels = data_dict.pop("spam")
spam_nrow, spam_ncol = spam_features.shape
spam_mean = spam_features.mean().to_numpy().reshape(1, spam_ncol)
spam_std = spam_features.std().to_numpy().reshape(1, spam_ncol)
spam_scaled = (spam_features - spam_mean)/spam_std
data_dict["spam_scaled"] = (spam_scaled, spam_labels)
{data_name:X.shape for data_name, (X,y) in data_dict.items()}

class Node:
    def __repr__(self):
        return "%s%s"%(self.__class__.__name__, self.value.shape)
```

```

class InitialNode(Node):
    def __init__(self, value):
        self.value = value
    def backward(self):
        pass

class Operation(Node):
    def backward(self):
        gradients = self.gradient()
        for parent_node, grad in zip(self.parents, gradients):
            if grad is not None and parent_node.value.shape != grad.shape:
                raise ValueError(
                    "value%s not same shape as grad%s"%(
                        str(parent_node.value.shape),
                        str(grad.shape)))
            parent_node.grad = grad
            parent_node.backward()

class mm(Operation):
    def __init__(self, feature_node, weight_node):
        self.parents = [feature_node, weight_node]
        self.value = np.matmul(feature_node.value, weight_node.value)
    def gradient(self):
        feature_node, weight_node = self.parents
        return[
            np.matmul(self.grad, weight_node.value.T),
            np.matmul(feature_node.value.T, self.grad)]

class logistic_loss(Operation):
    def __init__(self, pred_node, output_node):
        self.parents = [pred_node, output_node]
        output_vec = output_node.value
        if not ((output_vec == 1) | (output_vec == -1)).all():
            raise ValueError("Labels should be only -1 or 1")
        self.value = np.log(1 + np.exp(-output_vec * pred_node.value))

    def gradient(self):
        pred_node, output_node = self.parents
        # features X is b x p
        # weights W is p x u = 1
        # pred A is b x u = 1
        # where b is batch size
        # p is number of input features
        # u is number of outputs
        # grad_A(b x u) W(u x p)

```

```

pred_grad = -output_node.value/(
    1 + np.exp(
        output_node.value*
        pred_node.value
    )
)

```

```

return [pred_grad, None]

```

```

class CSV(torch.utils.data.Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels
    def __getitem__(self, item):
        return self.features[item,:], self.labels[item]
    def __len__(self):
        return len(self.labels)

```

```

class AutoMLP:
    def __init__(self, max_epochs, batch_size, step_size, units_per_layer):
        self.units_per_layer = units_per_layer
        self.max_epochs = max_epochs
        self.batch_size = batch_size
        self.step_size = step_size
        self.weight_node = InitialNode(
            np.repeat(0.0, self.units_per_layer[0]).reshape(self.units_per_layer[0], 1))

    def get_pred_node(self, batch_features):
        feature_node = InitialNode(np.array(batch_features))
        pred_node = mm(feature_node, self.weight_node)
        return pred_node

    def take_step(self, batch_features, batch_labels):
        label_node = InitialNode(np.array(batch_labels))
        pred_vec = self.get_pred_node(batch_features)
        loss_node = logistic_loss(pred_vec, label_node)
        loss_node.backward()
        gradient = self.weight_node.grad
        self.weight_node.value -= gradient * self.step_size
        return loss_node.value.mean()

    def fit(self, train_features, test_features):
        ds = CSV(train_features, test_features)
        dl = torch.utils.data.DataLoader(
            ds, batch_size = self.batch_size, shuffle = True)
        train_df_list = []
        for batch_features, batch_labels in dl:

```

```

        loss_value = self.take_step(batch_features, batch_labels)

def decision_function(self, X):
    pred_vec = self.get_pred_node(X)
    return pred_vec.value.reshape(len(pred_vec.value),)

def predict(self, X):
    pred_scores = self.decision_function(X)
    return np.where(pred_scores > 0, 1, 0)

class AutoGradLearnerCV:
    def __init__(self, max_epochs, batch_size, step_size, units_per_layer, n_splits):
        self.units_per_layer = units_per_layer
        self.max_epochs = max_epochs
        self.step_size = step_size
        self.batch_size = batch_size
        self.n_splits = n_splits

    def fit(self, train_features, train_labels):
        best_model = None
        train_nrow, train_ncol = train_features.shape
        times_to_repeat = int(math.ceil(train_nrow/self.n_splits))
        fold_id_vec = np.tile(np.arange(self.n_splits), times_to_repeat)[:train_nrow]
        np.random.shuffle(fold_id_vec)
        cv_data_list = []
        for epoch in range(1, self.max_epochs + 1):
            for validation_fold in range(self.n_splits):
                is_split = {
                    "subtrain": fold_id_vec != validation_fold,
                    "validation": fold_id_vec == validation_fold
                }
                split_data_dict = {}
                for set_name, is_set in is_split.items():
                    set_y = np.where(train_labels == 1, 1, -1).reshape(train_nrow, 1)
                    split_data_dict[set_name] = {
                        "n": len(set_y),
                        "X": train_features[is_set, :],
                        "y": set_y[is_set]}
                learner = AutoMLP(self.max_epochs, self.batch_size, self.step_size,
self.units_per_layer)
                learner.fit(split_data_dict["subtrain"]["X"], split_data_dict["subtrain"]["y"])
                for set_name, set_data in split_data_dict.items():
                    set_loss_value = learner.take_step(set_data["X"], set_data["y"])
                    cv_data_list.append(pd.DataFrame({
                        "set_name": [set_name],
                        "loss": float(set_loss_value),
                        "epoch": [epoch]

```

```

    )))

    self.cv_data = pd.concat(cv_data_list)
    best_epoch = self.cv_data.groupby('epoch')['loss'].mean().idxmin()
    best_learner = AutoMLP(best_epoch, self.batch_size, self.step_size, self.units_per_layer)
    best_learner.fit(train_features, np.where(train_labels == 1, 1, -1).reshape(train_nrow, 1))
    self.best_model = best_learner
    return self.cv_data

def predict(self, test_features):
    return self.best_model.predict(test_features)

```

```

accuracy_data_frames = []
loss_data_dict = {}
min_df_dict = {}
for data_name, (data_features, data_labels) in data_dict.items():
    kf = KFold(n_splits=3, shuffle=True, random_state=3)
    enum_obj = enumerate(kf.split(data_features))
    for fold_num, index_tup in enum_obj:
        zip_obj = zip(["train", "test"], index_tup)
        split_data = {}
        for set_name, set_indices in zip_obj:
            split_data[set_name] = (data_features.iloc[set_indices, :].to_numpy(),
                                     np.ravel(data_labels.iloc[set_indices]))
        train_features, train_labels = split_data["train"]
        nrow, ncol = train_features.shape
        test_features, test_labels = split_data["test"]

```

```

#KNN Classifier
knn = KNeighborsClassifier()
hp_parameters = {"n_neighbors": list(range(1, 21))}
grid = GridSearchCV(knn, hp_parameters, cv=5)
grid.fit(train_features, train_labels)
best_n_neighbors = grid.best_params_['n_neighbors']
print("Best N-Neighbors = ", best_n_neighbors)
knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(train_features, train_labels)
knn_pred = knn.predict(test_features)

```

```

# Logistic Regression
pipe = make_pipeline(StandardScaler(), LogisticRegressionCV(cv=3, max_iter=2000))
pipe.fit(train_features, train_labels)
lr_pred = pipe.predict(test_features)

```

```

#Featureless
y_train_series = pd.Series(train_labels)
most_frequent_class = y_train_series.value_counts().idxmax()
print("Most Frequent Class = ", most_frequent_class)

```

```

# create a featureless baseline
featureless_pred = np.repeat(most_frequent_class, len(test_features))

#AutoGradLearnerCV
model_units = {
    "linear": (ncol, 1),
    "deep": (ncol, 100, 10, 1)
}

#AutoGradLearnerCV_linear
linear_learner = AutoGradLearnerCV(50, 10, 0.015, [ncol, 1], 3)
linear_loss = linear_learner.fit(train_features, train_labels)
ll_pred = linear_learner.predict(test_features)

#AutoGradLearnerCV_deep
deep_learner = AutoGradLearnerCV(50, 10, 0.01, [ncol, 100, 10, 1], 3)
deep_loss = deep_learner.fit(train_features, train_labels)
dl_pred = deep_learner.predict(test_features)

linear_loss = linear_loss.groupby(['set_name', 'epoch']).mean().reset_index()
deep_loss = deep_loss.groupby(['set_name', 'epoch']).mean().reset_index()

valid_df = linear_loss.query("set_name=='validation'")
index_min = valid_df["loss"].argmin()
min_df = valid_df.query("epoch==%s" % (index_min + 1))

valid_df_deep = deep_loss.query("set_name=='validation'")
index_min_deep = valid_df_deep["loss"].argmin()
min_df_deep = valid_df_deep.query("epoch==%s" % (index_min_deep + 1))

min_df_dict[data_name] = {'min_df linear': min_df,
                          'min_df deep': min_df_deep}

loss_data_dict[data_name] = {'AutoGradLearnerCV Linear': linear_loss,
                              'AutoGradLearnerCV Deep': deep_loss}

# store predict data in dict
pred_dict = {'gridSearch + nearest neighbors': knn_pred,
             'linear_model': lr_pred,
             'AutoGradLearnerCV Linear': ll_pred,
             'AutoGradLearnerCV Deep': dl_pred,
             'featureless': featureless_pred}
test_accuracy = {}

for algorithm, predictions in pred_dict.items():
    #print(f"{algorithm}:", predictions.shape)

```

```

accuracy = np.mean(test_labels == predictions)
test_accuracy[algorithm] = accuracy

for algorithm, accuracy in test_accuracy.items():
    print(f"{algorithm} Test Accuracy: {accuracy * 100}")
    accuracy_df = pd.DataFrame({
        "data_set": [data_name],
        "fold_id": [fold_num],
        "algorithm": [algorithm],
        "accuracy": [test_accuracy[algorithm]]})
    accuracy_data_frames.append(accuracy_df)
    print(f"*****End of
{data_name}{{fold_num}}*****")

total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
print(total_accuracy_df)

zip_loss = loss_data_dict["zip"]
spam_loss = loss_data_dict["spam_scaled"]
zip_min = min_df_dict["zip"]
spam_min = min_df_dict["spam_scaled"]

gg = p9.ggplot() +\
    p9.geom_line(
        p9.aes(
            x = "epoch",
            y = "loss",
            color = "set_name"
        ),
        data = zip_loss["AutoGradLearnerCV Linear"]) +\
    p9.geom_point(
        p9.aes(
            x = "epoch",
            y = "loss",
            color = "set_name"
        ),
        data = zip_min["min_df linear"]) +\
    p9.ggtitle("Subtrain/Validation Loss vs Epochs(Zip Data - Linear)")

gg1 = p9.ggplot() +\
    p9.geom_line(
        p9.aes(
            x = "epoch",
            y = "loss",
            color = "set_name"
        ),
        data = zip_loss["AutoGradLearnerCV Deep"]) +\
    p9.geom_point(

```

```

p9.aes(
  x = "epoch",
  y = "loss",
  color = "set_name"
),
data = zip_min["min_df deep"]) +\
p9.ggtitle("Subtrain/Validation Loss vs Epochs(Zip Data - Deep)")

gg2 = p9.ggplot() +\
p9.geom_line(
  p9.aes(
    x = "epoch",
    y = "loss",
    color = "set_name"
  ),
  data = spam_loss["AutoGradLearnerCV Linear"]) +\
p9.geom_point(
  p9.aes(
    x = "epoch",
    y = "loss",
    color = "set_name"
  ),
  data = spam_min["min_df linear"]) +\
p9.ggtitle("Subtrain/Validation Loss vs Epochs(Spam_scaled Data - Linear)")

gg3 = p9.ggplot() +\
p9.geom_line(
  p9.aes(
    x = "epoch",
    y = "loss",
    color = "set_name"
  ),
  data = spam_loss["AutoGradLearnerCV Deep"]) +\
p9.geom_point(
  p9.aes(
    x = "epoch",
    y = "loss",
    color = "set_name"
  ),
  data = spam_min["min_df deep"]) +\
p9.ggtitle("Subtrain/Validation Loss vs Epochs(Spam_scaled Data - Deep)")

gg4 = p9.ggplot(total_accuracy_df, p9.aes(x = 'accuracy', y = 'algorithm')) +\
  p9.facet_grid('~data_set') + p9.geom_point()

gg.save("Zip_linear_SV_graph.png", height = 8, width = 12)
gg1.save("Zip_deep_SV_graph.png", height = 8, width = 12)

```



```

gg2.save("Spam_linear_SV_graph.png", height = 8, width = 12)
gg3.save("Spam_deep_SV_graph.png", height = 8, width = 12)
gg4.save("Accuracy_graph.png", height = 8, width = 12)

```

2. Output:

```

>>> for data_name, (data_features, data_labels) in data_dict.items():
...     kf = KFold(n_splits=3, shuffle=True, random_state=3)
...     enum_obj = enumerate(kf.split(data_features))
...     for fold_num, index_tup in enum_obj:
...         zip_obj = zip(["train", "test"], index_tup)
...         split_data = {}
...         for set_name, set_indices in zip_obj:
...             split_data[set_name] = (data_features.iloc[set_indices, :].to_numpy(),
...                                     np.ravel(data_labels.iloc[set_indices]))
...         train_features, train_labels = split_data["train"]
...     ...
...
...     for algorithm, accuracy in test_accuracy.items():
...         print(f"{algorithm} Test Accuracy: {accuracy * 100}")
...         accuracy_df = pd.DataFrame({
...             "data_set": [data_name],
...             "fold_id": [fold_num],
...             "algorithm": [algorithm],
...             "accuracy": [test_accuracy[algorithm]]})
...         accuracy_data_frames.append(accuracy_df)
...     print(f"*****End of
{data_name}{{fold_num}}*****")

```

```

Best N-Neighbors = 1
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 100.0
linear_model Test Accuracy: 99.51923076923077
AutoGradLearnerCV Linear Test Accuracy: 98.07692307692307
AutoGradLearnerCV Deep Test Accuracy: 99.51923076923077
featureless Test Accuracy: 58.65384615384615
*****End of zip(0)*****

```

```

Best N-Neighbors = 1
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 99.51923076923077
linear_model Test Accuracy: 99.03846153846155
AutoGradLearnerCV Linear Test Accuracy: 98.5576923076923
AutoGradLearnerCV Deep Test Accuracy: 98.5576923076923
featureless Test Accuracy: 57.21153846153846
*****End of zip(1)*****

```

```

Best N-Neighbors = 3
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 99.03381642512076

```

```

linear_model Test Accuracy: 99.03381642512076
AutoGradLearnerCV Linear Test Accuracy: 98.55072463768117
AutoGradLearnerCV Deep Test Accuracy: 98.55072463768117
featureless Test Accuracy: 57.00483091787439
*****End of zip(2)*****

Best N-Neighbors = 5
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 90.28683181225554
linear_model Test Accuracy: 91.52542372881356
AutoGradLearnerCV Linear Test Accuracy: 91.13428943937419
AutoGradLearnerCV Deep Test Accuracy: 91.59061277705347
featureless Test Accuracy: 60.88657105606258
*****End of spam_scaled(0)*****

Best N-Neighbors = 6
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 89.4393741851369
linear_model Test Accuracy: 91.78617992177314
AutoGradLearnerCV Linear Test Accuracy: 91.26466753585397
AutoGradLearnerCV Deep Test Accuracy: 91.39504563233378
featureless Test Accuracy: 60.104302477183836
*****End of spam_scaled(1)*****

Best N-Neighbors = 5
Most Frequent Class = 0
gridSearch + nearest neighbors Test Accuracy: 90.01956947162427
linear_model Test Accuracy: 92.49836921069797
AutoGradLearnerCV Linear Test Accuracy: 92.4331376386171
AutoGradLearnerCV Deep Test Accuracy: 91.71559034572732
featureless Test Accuracy: 60.79582517938682
*****End of spam_scaled(2)*****

```

```
>>> total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
```


```
>>> print(total_accuracy_df)
```

	data_set	fold_id	algorithm	accuracy
0	zip	0	gridSearch + nearest neighbors	1.000000
1	zip	0	linear_model	0.995192
2	zip	0	AutoGradLearnerCV Linear	0.980769
3	zip	0	AutoGradLearnerCV Deep	0.995192
4	zip	0	featureless	0.586538
5	zip	1	gridSearch + nearest neighbors	0.995192
6	zip	1	linear_model	0.990385
7	zip	1	AutoGradLearnerCV Linear	0.985577
8	zip	1	AutoGradLearnerCV Deep	0.985577
9	zip	1	featureless	0.572115
10	zip	2	gridSearch + nearest neighbors	0.990338
11	zip	2	linear_model	0.990338
12	zip	2	AutoGradLearnerCV Linear	0.985507

13	zip	2	AutoGradLearnerCV Deep	0.985507
14	zip	2	featureless	0.570048
15	spam_scaled	0	gridSearch + nearest neighbors	0.902868
16	spam_scaled	0	linear_model	0.915254
17	spam_scaled	0	AutoGradLearnerCV Linear	0.911343
18	spam_scaled	0	AutoGradLearnerCV Deep	0.915906
19	spam_scaled	0	featureless	0.608866
20	spam_scaled	1	gridSearch + nearest neighbors	0.894394
21	spam_scaled	1	linear_model	0.917862
22	spam_scaled	1	AutoGradLearnerCV Linear	0.912647
23	spam_scaled	1	AutoGradLearnerCV Deep	0.913950
24	spam_scaled	1	featureless	0.601043
25	spam_scaled	2	gridSearch + nearest neighbors	0.900196
26	spam_scaled	2	linear_model	0.924984
27	spam_scaled	2	AutoGradLearnerCV Linear	0.924331
28	spam_scaled	2	AutoGradLearnerCV Deep	0.917156
29	spam_scaled	2	featureless	0.607958

Accuracy Graph:

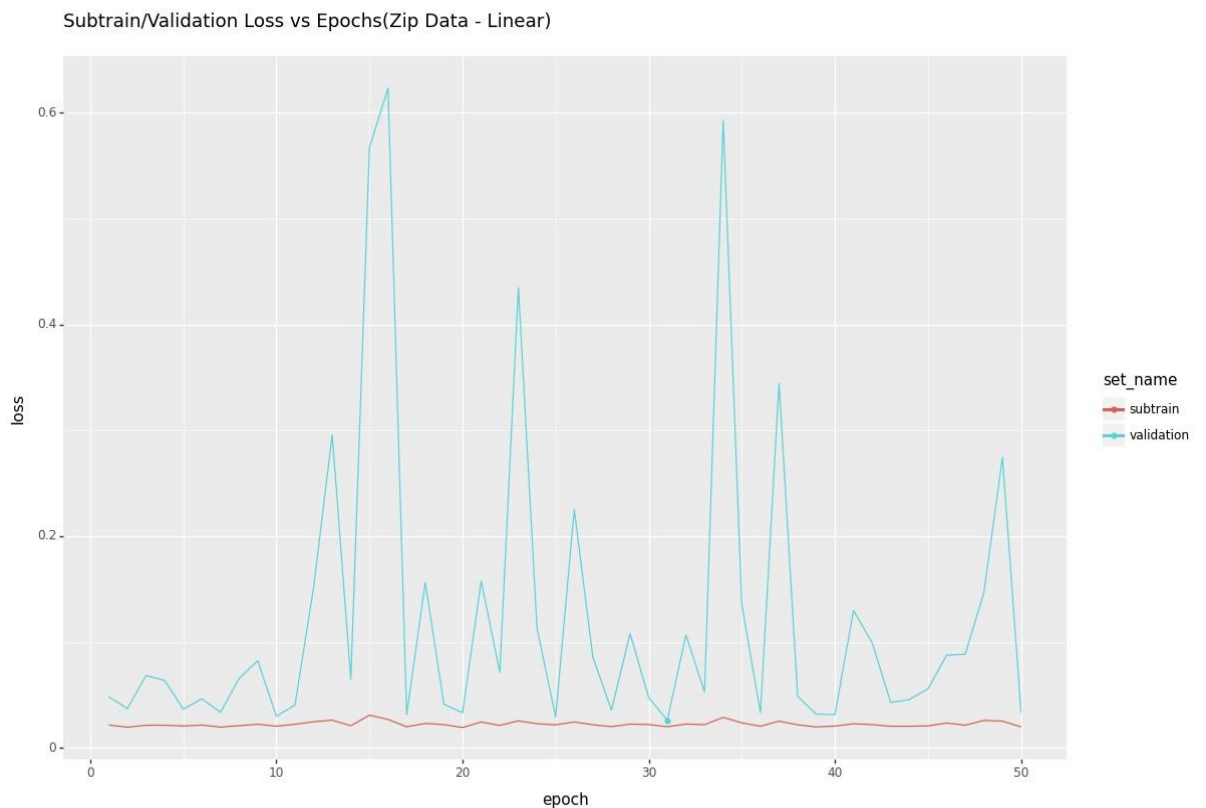
```
>>> gg4 = p9.ggplot(total_accuracy_df, p9.aes(x='accuracy', y='algorithm'))+\
...     p9.facet_grid('~data_set') + p9.geom_point()
>>> gg4.save("Accuracy_graph.png", height = 8, width = 12)
```



Linear subtrain/validation Loss graph (Zip):

```
>>> gg = p9.ggplot() +\  
...   p9.geom_line(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = zip_loss["AutoGradLearnerCV Linear"]) +\  
...   p9.geom_point(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = zip_min["min_df linear"]) +\  
...   p9.ggtitle("Subtrain/Validation Loss vs Epochs(Zip Data - Linear)")
```


```
>>> gg.save("Zip_linear_SV_graph.png", height = 8, width = 12)
```



Subtrain/Validation Loss Graph (Zip Data – Deep Model):

```
>>> gg1 = p9.ggplot() +\  
...   p9.geom_line(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = zip_loss["AutoGradLearnerCV Deep"]) +\  
...   p9.geom_point(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = zip_min["min_df deep"]) +\  
...   p9.ggtitle("Subtrain/Validation Loss vs Epochs(Zip Data - Deep)")
```


Subtrain/Validation Loss vs Epochs(Zip Data - Deep)



Subtrain/Validation Loss Graph (Spam_scaled Data – Linear Model):

```
>>> gg2 = p9.ggplot() +\  
...   p9.geom_line(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = spam_loss["AutoGradLearnerCV Linear"]) +\  
...   p9.geom_point(  
...     p9.aes(  
...       x = "epoch",  
...       y = "loss",  
...       color = "set_name"  
...     ),  
...     data = spam_min["min_df linear"]) +\  
...   p9.ggtitle("Subtrain/Validation Loss vs Epochs(Spam_scaled Data - Linear)")
```

Subtrain/Validation Loss vs Epochs(Spam_scaled Data - Linear)



Subtrain/Validation Loss Graph (Spam_scaled Data – Deep Model):

