

CS599 (Deep Learning)

Homework – 09

1. Python Code:

```
import torch
import pandas as pd
import matplotlib
import numpy as np
import math
matplotlib.use("agg")

from sklearn.model_selection import KFold, GridSearchCV, ParameterGrid
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LassoCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
from collections import Counter

data_set_dict = {"forest_fires": ("forestfires.csv", ",", True),
                 "air_foil": ("airfoil_self_noise.tsv", "\t", False)}
data_dict = {}
hist_df_list = []

for data_name, (file_name, sep, log_trans) in data_set_dict.items():
    data_df = pd.read_csv(file_name, sep = sep, header = 0)
    data_nrow, data_ncol = data_df.shape
    label_col_num = data_ncol - 1
    data_label_vec = data_df.iloc[:, label_col_num]
    if log_trans:
        data_label_vec = np.log(data_label_vec + 1)
    label_sd = math.sqrt(data_label_vec.var())
    standard_label_vec = (
        data_label_vec - data_label_vec.mean()
    ) / label_sd
    is_feature_col = (
        np.arange(data_ncol) != label_col_num
    ) & (
        data_df.dtypes != "object"
    )
    data_features = data_df.loc[:, is_feature_col]
    feature_nrow, feature_ncol = data_features.shape
```

```

feature_mean = data_features.mean().to_numpy().reshape(1, feature_ncol)
feature_std = data_features.std().to_numpy().reshape(1, feature_ncol)
feature_scaled = (data_features - feature_mean)/feature_std
print("%s %s" %(data_name, data_features.shape))
input_tensor = torch.from_numpy(
    feature_scaled.to_numpy()
).float()
output_tensor = torch.from_numpy(
    standard_label_vec.to_numpy()
).float().reshape(data_nrow, 1)
data_dict[data_name] = (feature_scaled, standard_label_vec)
hist_df_list.append(pd.DataFrame({
    "data_name": data_name,
    "label": standard_label_vec})
)
hist_df = pd.concat(hist_df_list)

```

```

class TorchModel(torch.nn.Module):
    def __init__(self, units_per_layer):
        super(TorchModel, self).__init__()
        seq_args = []
        second_to_last = len(units_per_layer)-1
        for layer_i in range(second_to_last):
            next_i = layer_i+1
            layer_units = units_per_layer[layer_i]
            next_units = units_per_layer[next_i]
            seq_args.append(torch.nn.Linear(layer_units, next_units))
            if layer_i < second_to_last-1:
                seq_args.append(torch.nn.ReLU())
        self.stack = torch.nn.Sequential(*seq_args)
    def forward(self, features):
        return self.stack(features)

```

```

class CSV(torch.utils.data.Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels
    def __getitem__(self, item):
        return self.features[item,:], self.labels[item]
    def __len__(self):
        return len(self.labels)

```

```

class TorchLearner:
    def __init__(
        self, units_per_layer, step_size=0.1,
        batch_size=20, max_epochs=50):
        self.max_epochs = max_epochs
        self.batch_size=batch_size

```

```

self.model = TorchModel(units_per_layer)
self.loss_fun = torch.nn.MSELoss()
self.optimizer = torch.optim.SGD(
    self.model.parameters(), lr=step_size)
def fit(self, split_data_dict):
    ds = CSV(
        split_data_dict["subtrain"]["X"],
        split_data_dict["subtrain"]["y"])
    dl = torch.utils.data.DataLoader(
        ds, batch_size=self.batch_size, shuffle=True)
    train_df_list = []
    for epoch_number in range(self.max_epochs):
        #print(epoch_number)
        for batch_features, batch_labels in dl:
            self.optimizer.zero_grad()
            loss_value = self.loss_fun(
                self.model(batch_features), batch_labels)
            loss_value.backward()
            self.optimizer.step()
        for set_name, set_data in split_data_dict.items():
            pred_vec = self.model(set_data["X"])
            set_loss_value = self.loss_fun(pred_vec, set_data["y"])
            train_df_list.append(pd.DataFrame({
                "set_name": [set_name],
                "loss": float(set_loss_value),
                "epoch": [epoch_number]
            })))
    self.train_df = pd.concat(train_df_list)
def decision_function(self, test_features):
    with torch.no_grad():
        pred_vec = self.model(test_features)
    return pred_vec.numpy()

def predict(self, test_features):
    pred_scores = self.decision_function(test_features)
    return pred_scores

```

```

class TorchLearnerCV:

```

```

    def __init__(self, n_folds, units_per_layer=[data_ncol,1]):
        self.units_per_layer = units_per_layer
        self.n_folds = n_folds
    def fit(self, train_features, train_labels):
        train_nrow, train_ncol = train_features.shape
        times_to_repeat=int(math.ceil(train_nrow/self.n_folds))
        fold_id_vec = np.tile(torch.arange(self.n_folds), times_to_repeat)[:train_nrow]
        np.random.shuffle(fold_id_vec)
        cv_data_list = []

```

```

for validation_fold in range(self.n_folds):
    is_split = {
        "subtrain":fold_id_vec != validation_fold,
        "validation":fold_id_vec == validation_fold
    }
    split_data_dict = {}
    for set_name, is_set in is_split.items():
        set_y = train_labels[is_set]
        split_data_dict[set_name] = {
            "X":train_features[is_set,:],
            "y":set_y}
    learner = TorchLearner(self.units_per_layer)
    learner.fit(split_data_dict)
    cv_data_list.append(learner.train_df)
self.cv_data = pd.concat(cv_data_list)
self.train_df = self.cv_data.groupby(["set_name","epoch"]).mean().reset_index()
#print(self.train_df)
valid_df = self.train_df.query("set_name=='validation'")
#print(valid_df)
best_epochs = valid_df["loss"].argmin()
self.min_df = valid_df.query("epoch==%s"%(best_epochs))
print("Best Epoch: ", best_epochs)
self.final_learner = TorchLearner(self.units_per_layer, max_epochs=(best_epochs + 1))
self.final_learner.fit({"subtrain":{"X":train_features,"y":train_labels}})
return self.cv_data
def predict(self, test_features):
    return self.final_learner.predict(test_features)

```

```

test_loss_data_frames = []
loss_data_dict = {}
min_df_dict = {}
for data_name, (data_features, data_labels) in data_dict.items():
    kf = KFold(n_splits=3, shuffle=True, random_state=3)
    enum_obj = enumerate(kf.split(data_features))
    for fold_num, index_tup in enum_obj:
        zip_obj = zip(["train", "test"], index_tup)
        split_data = {}
        for set_name, set_indices in zip_obj:
            split_data[set_name] = (torch.from_numpy(data_features.iloc[set_indices,
:].to_numpy()).float(),
                                torch.from_numpy(np.ravel(data_labels.iloc[set_indices])).float())
        #x = {data_name:X.shape for data_name, (X,y) in split_data.items()}
        #print(f"{data_name}: ", x)
        train_features, train_labels = split_data["train"]
        nrow, ncol = train_features.shape
        #print(f"{data_name}: ", nrow, ncol)

```

```

test_features, test_labels = split_data["test"]

#kneighbors
knn = KNeighborsRegressor()
hp_parameters = {"n_neighbors": list(range(1, 21))}
grid = GridSearchCV(knn, hp_parameters, cv=3)
grid.fit(train_features, train_labels)
best_n_neighbors = grid.best_params_['n_neighbors']
print("Best N-Neighbors = ", best_n_neighbors)
knn = KNeighborsRegressor(n_neighbors=best_n_neighbors)
knn.fit(train_features, train_labels)
knn_pred = knn.predict(test_features)
#print(test_labels)
#print(knn_pred)
#loss = mean_squared_error(test_labels, knn_pred)
#print(f"Knn Loss {data_name} : ", loss)

#linear model
pipe = make_pipeline(StandardScaler(), LassoCV(cv=3, max_iter=2000))
pipe.fit(train_features, train_labels)
lr_pred = pipe.predict(test_features)

#loss_linear = mean_squared_error(test_labels, lr_pred)
#print(f"Linear_loss {data_name} : ", loss_linear)

#Featureless
y_train_series = pd.Series(train_labels)
mean_train_label = y_train_series.mean()
print("Mean Train Label = ", mean_train_label)

# create a featureless baseline
featureless_pred = np.repeat(mean_train_label, len(test_features))
#featureless_loss = mean_squared_error(test_labels, featureless_pred)
#print(f"Featureless Loss {data_name} : ", featureless_loss)

#TorchLearnerCV
linear_learner = TorchLearnerCV(3, [ncol, 1])
#print("ncol:", ncol)
linear_loss = linear_learner.fit(train_features, train_labels.reshape(nrow,1))
ll_pred = linear_learner.predict(test_features)
#print(ll_pred)
#loss_torchlinear = mean_squared_error(test_labels, ll_pred)
#print(f"Torch Linear_loss {data_name} : ", loss_torchlinear)

#TorchLearnerCV + Deep
deep_learner = TorchLearnerCV(3, [ncol, 100, 10, 1])
deep_loss = deep_learner.fit(train_features, train_labels.reshape(nrow, 1))
dl_pred = deep_learner.predict(test_features)

```

```

#loss_deeplearner = mean_squared_error(test_labels, dl_pred)
#print(f"Torch Deep_loss {data_name} : ", loss_deeplearner)

linear_loss = linear_loss.groupby(['set_name', 'epoch']).mean().reset_index()
deep_loss = deep_loss.groupby(['set_name', 'epoch']).mean().reset_index()

valid_df = linear_loss.query("set_name=='validation'")
index_min = valid_df["loss"].argmin()
min_df = valid_df.query("epoch==%s" % index_min)

valid_df_deep = deep_loss.query("set_name=='validation'")
index_min_deep = valid_df_deep["loss"].argmin()
min_df_deep = valid_df_deep.query("epoch==%s" % index_min_deep)

min_df_dict[data_name] = {'min_df linear': min_df,
                          'min_df deep': min_df_deep}

loss_data_dict[data_name] = {'TorchLearnerCV Linear': linear_loss,
                             'TorchLearnerCV Deep': deep_loss}

# store predict data in dict
pred_dict = {'KNeighborsRegressor + GridSearchCV': knn_pred,
             'LassoCV': lr_pred,
             'TorchLearnerCV Linear': ll_pred,
             'TorchLearnerCV Deep': dl_pred,
             'featureless': featureless_pred}
test_square_loss = {}
for algorithm, predictions in pred_dict.items():
    #print(f"{algorithm}:", predictions.shape)
    test_loss = mean_squared_error(test_labels, predictions)
    #accuracy = np.mean(test_labels == predictions)
    test_square_loss[algorithm] = test_loss

for algorithm, test_loss in test_square_loss.items():
    print(f"{algorithm} Test Square Loss: {test_loss}")
    test_loss_df = pd.DataFrame({
        "data_set": [data_name],
        "fold_id": [fold_num],
        "algorithm": [algorithm],
        "test square loss": [test_square_loss[algorithm]]})
    test_loss_data_frames.append(test_loss_df)
    print(f"*****End of
{data_name}{{fold_num}}*****")

total_test_loss_df = pd.concat(test_loss_data_frames, ignore_index = True)
print(total_test_loss_df)

```

```

import plotnine as p9
gg = p9.ggplot(total_test_loss_df, p9.aes(x='test square loss', y='algorithm'))+\
    p9.facet_grid('~data_set') + p9.geom_point()

gg.save("Test_square_loss.png", height = 8, width = 12)

forest_fires_loss = loss_data_dict["forest_fires"]
air_foil_loss = loss_data_dict["air_foil"]
forest_fires_min = min_df_dict["forest_fires"]
air_foil_min = min_df_dict["air_foil"]

gg1 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
forest_fires_loss["TorchLearnerCV Linear"])\
    + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
forest_fires_min["min_df linear"]) #p9.geom_line(p9.aes(fill='setname'))

gg1.save("Torch_validation_graph1.png", height = 8, width = 12)

gg2 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
forest_fires_loss["TorchLearnerCV Deep"])\
    + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
forest_fires_min["min_df deep"]) #p9.geom_line(p9.aes(fill='setname'))

gg2.save("Torch_validation_graph2.png", height = 8, width = 12)

gg3 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_loss["TorchLearnerCV Linear"])\
    + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_min["min_df linear"]) #p9.geom_line(p9.aes(fill='setname'))

gg3.save("Torch_validation_graph3.png", height = 8, width = 12)

gg4 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_loss["TorchLearnerCV Deep"])\
    + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_min["min_df deep"]) #p9.geom_line(p9.aes(fill='setname'))

gg4.save("Torch_validation_graph4.png", height = 8, width = 12)

```

2. Output:

```

>>> for data_name, (data_features, data_labels) in data_dict.items():
...     kf = KFold(n_splits=3, shuffle=True, random_state=3)
...     enum_obj = enumerate(kf.split(data_features))
...     for fold_num, index_tup in enum_obj:

```

```

...     zip_obj = zip(["train", "test"], index_tup)
...     split_data = {}
...     for set_name, set_indices in zip_obj:
...         split_data[set_name] = (torch.from_numpy(data_features.iloc[set_indices,
:]
.to_numpy()).float(),
...         torch.from_numpy(np.ravel(data_labels.iloc[set_indices])).float())
...     #x = {data_name:X.shape for data_name, (X,y) in split_data.items()}
...     ...
...
...     for algorithm, test_loss in test_square_loss.items():
...         print(f"{algorithm} Test Square Loss: {test_loss}")
...         test_loss_df = pd.DataFrame({
...             "data_set": [data_name],
...             "fold_id": [fold_num],
...             "algorithm": [algorithm],
...             "test square loss": [test_square_loss[algorithm]]})
...         test_loss_data_frames.append(test_loss_df)
...         print(f"*****End of
{data_name}{{fold_num}}*****")

```

Best N-Neighbors = 19

Mean Train Label = 0.0049645514

Best Epoch: 0

Best Epoch: 2

KNeighborsRegressor + GridSearchCV Test Square Loss: 1.0609387159347534

LassoCV Test Square Loss: 1.0739297224352429

TorchLearnerCV Linear Test Square Loss: 1.2798919677734375

TorchLearnerCV Deep Test Square Loss: 1.073811411857605

featureless Test Square Loss: 1.0739296674728394

*****End of forest_fires(0)*****

Best N-Neighbors = 20

Mean Train Label = -0.019172536

Best Epoch: 0

Best Epoch: 3

KNeighborsRegressor + GridSearchCV Test Square Loss: 0.988243818283081

LassoCV Test Square Loss: 0.9922808202492903

TorchLearnerCV Linear Test Square Loss: 1.0052711963653564

TorchLearnerCV Deep Test Square Loss: 0.989805281162262

featureless Test Square Loss: 0.9922808408737183

*****End of forest_fires(1)*****

Best N-Neighbors = 17

Mean Train Label = 0.014222353

Best Epoch: 10

Best Epoch: 0

KNeighborsRegressor + GridSearchCV Test Square Loss: 1.0026485919952393

LassoCV Test Square Loss: 0.9305249905775753

TorchLearnerCV Linear Test Square Loss: 1.1973176002502441

TorchLearnerCV Deep Test Square Loss: 0.9264524579048157


```

featureless Test Square Loss: 0.9305248856544495
*****End of forest_fires(2)*****

Best N-Neighbors = 13
Mean Train Label = -0.034039237
Best Epoch: 42
Best Epoch: 44
KNeighborsRegressor + GridSearchCV Test Square Loss: 0.31644493341445923
LassoCV Test Square Loss: 0.47843713543411764
TorchLearnerCV Linear Test Square Loss: 0.5696210265159607
TorchLearnerCV Deep Test Square Loss: 0.1302846521139145
featureless Test Square Loss: 0.9787271022796631
*****End of air_foil(0)*****

Best N-Neighbors = 9
Mean Train Label = 0.00056366913
Best Epoch: 44
Best Epoch: 49
KNeighborsRegressor + GridSearchCV Test Square Loss: 0.2813766300678253
LassoCV Test Square Loss: 0.48135759164756325
TorchLearnerCV Linear Test Square Loss: 0.6965237259864807
TorchLearnerCV Deep Test Square Loss: 0.2059951275587082
featureless Test Square Loss: 0.9960943460464478
*****End of air_foil(1)*****

Best N-Neighbors = 11
Mean Train Label = 0.033475567
Best Epoch: 42
Best Epoch: 45
KNeighborsRegressor + GridSearchCV Test Square Loss: 0.3124180734157562
LassoCV Test Square Loss: 0.5237005567324898
TorchLearnerCV Linear Test Square Loss: 0.5518736839294434
TorchLearnerCV Deep Test Square Loss: 0.3650369644165039
featureless Test Square Loss: 1.0345804691314697
*****End of air_foil(2)*****

```

```
>>> total_test_loss_df = pd.concat(test_loss_data_frames, ignore_index = True)
```

```
>>> print(total_test_loss_df)
```


	data_set	fold_id	algorithm	test square loss
0	forest_fires	0	KNeighborsRegressor + GridSearchCV	1.060939
1	forest_fires	0	LassoCV	1.073930
2	forest_fires	0	TorchLearnerCV Linear	1.279892
3	forest_fires	0	TorchLearnerCV Deep	1.073811
4	forest_fires	0	featureless	1.073930
5	forest_fires	1	KNeighborsRegressor + GridSearchCV	0.988244
6	forest_fires	1	LassoCV	0.992281
7	forest_fires	1	TorchLearnerCV Linear	1.005271
8	forest_fires	1	TorchLearnerCV Deep	0.989805
9	forest_fires	1	featureless	0.992281

10	forest_fires	2	KNeighborsRegressor + GridSearchCV	1.002649
11	forest_fires	2	LassoCV	0.930525
12	forest_fires	2	TorchLearnerCV Linear	1.197318
13	forest_fires	2	TorchLearnerCV Deep	0.926452
14	forest_fires	2	featureless	0.930525
15	air_foil	0	KNeighborsRegressor + GridSearchCV	0.316445
16	air_foil	0	LassoCV	0.478437
17	air_foil	0	TorchLearnerCV Linear	0.569621
18	air_foil	0	TorchLearnerCV Deep	0.130285
19	air_foil	0	featureless	0.978727
20	air_foil	1	KNeighborsRegressor + GridSearchCV	0.281377
21	air_foil	1	LassoCV	0.481358
22	air_foil	1	TorchLearnerCV Linear	0.696524
23	air_foil	1	TorchLearnerCV Deep	0.205995
24	air_foil	1	featureless	0.996094
25	air_foil	2	KNeighborsRegressor + GridSearchCV	0.312418
26	air_foil	2	LassoCV	0.523701
27	air_foil	2	TorchLearnerCV Linear	0.551874
28	air_foil	2	TorchLearnerCV Deep	0.365037
29	air_foil	2	featureless	1.034580

Test Loss Square Graph:

```
>>> gg = p9.ggplot(total_test_loss_df, p9.aes(x = 'test square loss', y = 'algorithm'))+\
...     p9.facet_grid('~data_set') + p9.geom_point()
```


```
>>> gg.save("Test_square_loss.png", height = 8, width = 12)
```



Linear subtrain/validation loss graph (Forest Fires):

```
>>> gg1 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data
= forest_fires_loss["TorchLearnerCV Linear"])\
... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
forest_fires_min["min_df linear"]) #p9.geom_line(p9.aes(fill='setname'))


>>> gg1.save("Torch_validation_graph1.png", height = 8, width = 12)
```



Linear subtrain/validation loss graph (Air Foil):

```
>>> gg3 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data
= air_foil_loss["TorchLearnerCV Linear"])\
... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_min["min_df linear"]) #p9.geom_line(p9.aes(fill='setname'))


>>> gg3.save("Torch_validation_graph3.png", height = 8, width = 12)
```



Deep subtrain/validation loss graph (Forest Fires):

```
>>> gg2 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data = forest_fires_loss["TorchLearnerCV Deep"])\
... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = forest_fires_min["min_df deep"]) #p9.geom_line(p9.aes(fill='setname'))
```

```
>>> gg2.save("Torch_validation_graph2.png", height = 8, width = 12)
```



Deep subtrain/validation loss graph (Air Foil):

```
>>> gg4 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data
= air_foil_loss["TorchLearnerCV Deep"])\
... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
air_foil_min["min_df deep"]) #p9.geom_line(p9.aes(fill='setname'))

>>> gg4.save("Torch_validation_graph4.png", height = 8, width = 12)
```

