

## CS599 (Deep Learning)

### Homework – 10

#### 1. Python Code:

```
import torch
import pandas as pd
import matplotlib
matplotlib.use("agg")
import numpy as np
import plotnine as p9
import math
import torchvision

from sklearn.model_selection import KFold, GridSearchCV, ParameterGrid
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from collections import Counter

data_set_dict = {"zip": ("zip.test.gz", 0),
                 }
data_dict = {}

for data_name, (file_name, label_col_num) in data_set_dict.items():
    data_df = pd.read_csv(file_name, sep=" ", header=None)
    data_nrow, data_ncol = data_df.shape
    data_label_vec = data_df.iloc[:, label_col_num]
    is_label_col = data_df.columns == label_col_num
    data_features = data_df.iloc[:, ~is_label_col]
    data_labels = data_df.iloc[:, is_label_col]
    print("%s %s" %(data_name, data_features.shape))
    data_dict[data_name] = (
        torch.from_numpy(data_features.to_numpy()).float(),
        torch.from_numpy(data_labels.to_numpy()).flatten()
    )

ds = torchvision.datasets.MNIST(
    root="c:/Users/Anudeep Kumar/OneDrive/Desktop/Fall 2023/CS599-Deep
Learning/Homework/HW10",
    download=True,
    transform=torchvision.transforms.ToTensor(),
    train=False)
```

```

dl = torch.utils.data.DataLoader(ds, batch_size=len(ds), shuffle=False)
for mnist_features, mnist_labels in dl:
    pass
mnist_features.flatten(start_dim=1)
mnist_labels.numpy()
data_dict["MNIST"] = (mnist_features.flatten(start_dim=1), mnist_labels)

```

```

class TorchModel(torch.nn.Module):
    def __init__(self, units_per_layer):
        super(TorchModel, self).__init__()
        seq_args = []
        second_to_last = len(units_per_layer)-1
        for layer_i in range(second_to_last):
            next_i = layer_i+1
            layer_units = units_per_layer[layer_i]
            next_units = units_per_layer[next_i]
            seq_args.append(torch.nn.Linear(layer_units, next_units))
            if layer_i < second_to_last-1:
                seq_args.append(torch.nn.ReLU())
        self.stack = torch.nn.Sequential(*seq_args)
    def forward(self, features):
        return self.stack(features)

```

```

class CSV(torch.utils.data.Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels
    def __getitem__(self, item):
        return self.features[item,:], self.labels[item]
    def __len__(self):
        return len(self.labels)

```

```

class TorchLearner:
    def __init__(
        self, units_per_layer, step_size=0.1,
        batch_size=20, max_epochs=100):
        self.max_epochs = max_epochs
        self.batch_size=batch_size
        self.model = TorchModel(units_per_layer)
        self.loss_fun = torch.nn.CrossEntropyLoss()
        self.optimizer = torch.optim.SGD(
            self.model.parameters(), lr=step_size)
    def fit(self, split_data_dict):
        ds = CSV(
            split_data_dict["subtrain"]["X"],
            split_data_dict["subtrain"]["y"])
        dl = torch.utils.data.DataLoader(
            ds, batch_size=self.batch_size, shuffle=True)

```

```

train_df_list = []
for epoch_number in range(self.max_epochs):
    #print(epoch_number)
    for batch_features, batch_labels in dl:
        self.optimizer.zero_grad()
        loss_value = self.loss_fun(
            self.model(batch_features), batch_labels)
        loss_value.backward()
        self.optimizer.step()
    for set_name, set_data in split_data_dict.items():
        pred_vec = self.model(set_data["X"])
        set_loss_value = self.loss_fun(pred_vec, set_data["y"])
        train_df_list.append(pd.DataFrame({
            "set_name": [set_name],
            "loss": float(set_loss_value),
            "epoch": [epoch_number]
        })))
    self.train_df = pd.concat(train_df_list)
def decision_function(self, test_features):
    with torch.no_grad():
        pred_vec = self.model(test_features)
    return pred_vec

def predict(self, test_features):
    pred_scores = self.decision_function(test_features)
    _, predicted = torch.max(pred_scores, 1)
    return predicted.numpy()

class TorchLearnerCV:
    def __init__(self, n_folds, units_per_layer):
        self.units_per_layer = units_per_layer
        self.n_folds = n_folds
    def fit(self, train_features, train_labels):
        train_nrow, train_ncol = train_features.shape
        times_to_repeat = int(math.ceil(train_nrow/self.n_folds))
        fold_id_vec = np.tile(torch.arange(self.n_folds), times_to_repeat)[:train_nrow]
        np.random.shuffle(fold_id_vec)
        cv_data_list = []
        for validation_fold in range(self.n_folds):
            is_split = {
                "subtrain": fold_id_vec != validation_fold,
                "validation": fold_id_vec == validation_fold
            }
            split_data_dict = {}
            for set_name, is_set in is_split.items():
                set_y = train_labels[is_set]
                split_data_dict[set_name] = {
                    "X": train_features[is_set,:],

```

```

        "y":set_y}
    learner = TorchLearner(self.units_per_layer)
    learner.fit(split_data_dict)
    cv_data_list.append(learner.train_df)
    self.cv_data = pd.concat(cv_data_list)
    self.train_df = self.cv_data.groupby(["set_name", "epoch"]).mean().reset_index()
    #print(self.train_df)
    valid_df = self.train_df.query("set_name=='validation'")
    #print(valid_df)
    best_epochs = valid_df["loss"].argmin()
    self.min_df = valid_df.query("epoch==%s"%(best_epochs))
    print("Best Epoch: ", best_epochs)
    self.final_learner = TorchLearner(self.units_per_layer, max_epochs=(best_epochs + 1))
    self.final_learner.fit({"subtrain":{"X":train_features,"y":train_labels}})
    return self.cv_data
def predict(self, test_features):
    return self.final_learner.predict(test_features)

```

```

accuracy_data_frames = []
loss_data_dict = {}
min_df_dict = {}
for data_name, (data_features, data_labels) in data_dict.items():
    kf = KFold(n_splits=3, shuffle=True, random_state=3)
    enum_obj = enumerate(kf.split(data_features))
    for fold_num, index_tup in enum_obj:
        zip_obj = zip(["train", "test"], index_tup)
        split_data = {}
        for set_name, set_indices in zip_obj:
            split_data[set_name] = (data_features, data_labels)
        #x = {data_name:X.shape for data_name, (X,y) in split_data.items()}
        #print(f"{data_name}: ", x)
        train_features, train_labels = split_data["train"]
        nrow, ncol = train_features.shape
        print(f"{data_name}: ", nrow, ncol)
        test_features, test_labels = split_data["test"]

    #kneighbors
    knn = KNeighborsClassifier()
    hp_parameters = {"n_neighbors": list(range(1, 21))}
    grid = GridSearchCV(knn, hp_parameters, cv=3)
    grid.fit(train_features, train_labels)
    best_n_neighbors = grid.best_params_['n_neighbors']
    print("Best N-Neighbors = ", best_n_neighbors)
    knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
    knn.fit(train_features, train_labels)
    knn_pred = knn.predict(test_features)
    #print(knn_pred)
    #loss = mean_squared_error(test_labels, knn_pred)

```

```

#print(f"Knn Loss {data_name} : ", loss)

#linear model
pipe = make_pipeline(StandardScaler(), LogisticRegressionCV(cv=3, max_iter=2000))
pipe.fit(train_features, train_labels)
lr_pred = pipe.predict(test_features)
#print(lr_pred)
#loss_linear = mean_squared_error(test_labels, lr_pred)
#print(f"Linear_loss {data_name} : ", loss_linear)

#Featureless
y_train_series = pd.Series(train_labels)
#mean_train_label = y_train_series.mean()
#print("Mean Train Label = ", mean_train_label)

# create a featureless baseline
most_frequent_label = y_train_series.value_counts().idxmax()
print("Most Frequent Label = ", most_frequent_label)

featureless_pred = np.repeat(most_frequent_label, len(test_features))
#featureless_loss = mean_squared_error(test_labels, featureless_pred)
#print(f"Featureless Loss {data_name} : ", featureless_loss)

#TorchLearnerCV
linear_learner = TorchLearnerCV(3, [ncol, 10])
#print("ncol:", ncol)
linear_loss = linear_learner.fit(train_features, train_labels)
ll_pred = linear_learner.predict(test_features)
#print(ll_pred)
#loss_torchlinear = mean_squared_error(test_labels, ll_pred)
#print(f"Torch Linear_loss {data_name} : ", loss_torchlinear)

#TorchLearnerCV + Deep
deep_learner = TorchLearnerCV(3, [ncol, 100, 10, 10])
deep_loss = deep_learner.fit(train_features, train_labels)
dl_pred = deep_learner.predict(test_features)
#print(dl_pred)
#loss_deeplearner = mean_squared_error(test_labels, dl_pred)
#print(f"Torch Deep_loss {data_name} : ", loss_deeplearner)

linear_loss = linear_loss.groupby(['set_name', 'epoch']).mean().reset_index()
deep_loss = deep_loss.groupby(['set_name', 'epoch']).mean().reset_index()

valid_df = linear_loss.query("set_name=='validation'")
index_min = valid_df["loss"].argmin()
min_df = valid_df.query("epoch==%s" % index_min)

valid_df_deep = deep_loss.query("set_name=='validation'")

```

```

index_min_deep = valid_df_deep["loss"].argmin()
min_df_deep = valid_df_deep.query("epoch==%s" % index_min_deep)

min_df_dict[data_name] = {'min_df linear': min_df,
                          'min_df deep': min_df_deep}

loss_data_dict[data_name] = {'TorchLearnerCV Linear': linear_loss,
                              'TorchLearnerCV Deep': deep_loss}

# store predict data in dict
pred_dict = {'KNeighborsClassifier + GridSearchCV': knn_pred,
             'LogisticRegressionCV': lr_pred,
             'TorchLearnerCV Linear': ll_pred,
             'TorchLearnerCV Deep': dl_pred,
             'featureless': featureless_pred}
test_accuracy = {}
for algorithm, predictions in pred_dict.items():
    #print(f"{algorithm}:", predictions.shape)
    #test_loss = mean_squared_error(test_labels, predictions)
    accuracy = accuracy_score(test_labels, predictions)
    test_accuracy[algorithm] = accuracy

for algorithm, accuracy in test_accuracy.items():
    print(f"{algorithm} Test Accuracy: {accuracy * 100}")
    accuracy_df = pd.DataFrame({
        "data_set": [data_name],
        "fold_id": [fold_num],
        "algorithm": [algorithm],
        "accuracy": [test_accuracy[algorithm]]})
    accuracy_data_frames.append(accuracy_df)
    print(f"*****End of
{data_name}{{fold_num}}*****")

total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)

print(total_accuracy_df)

import plotnine as p9
gg = p9.ggplot(total_accuracy_df, p9.aes(x='accuracy', y = 'algorithm'))+\
    p9.facet_grid('~data_set') + p9.geom_point()

gg.save("c:/Users/Anudeep Kumar/OneDrive/Desktop/Fall 2023/CS599-Deep
Learning/Homework/HW10/output.png", height = 8, width = 12)

zip_loss = loss_data_dict["zip"]
mnist_loss = loss_data_dict["MNIST"]

```

```

zip_min = min_df_dict["zip"]
mnist_min = min_df_dict["MNIST"]

gg1 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
zip_loss["TorchLearnerCV Linear"])\
+ p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_min["min_df
linear"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(zip Data - Linear)")

gg1.save("Torch_validation_graph1.png", height = 8, width = 12)

gg2 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
zip_loss["TorchLearnerCV Deep"])\
+ p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_min["min_df
deep"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(zip Data - Deep)")

gg2.save("Torch_validation_graph2.png", height = 8, width = 12)

gg3 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
mnist_loss["TorchLearnerCV Linear"])\
+ p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
mnist_min["min_df linear"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(MNIST Data -
Linear)")

gg3.save("Torch_validation_graph3.png", height = 8, width = 12)

gg4 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data =
mnist_loss["TorchLearnerCV Deep"])\
+ p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data =
mnist_min["min_df deep"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(MNIST Data -
Deep)")

gg4.save("Torch_validation_graph4.png", height = 8, width = 12)

```

## 2. Output:

```

>>> for data_name, (data_features, data_labels) in data_dict.items():
...     kf = KFold(n_splits=3, shuffle=True, random_state=3)
...     enum_obj = enumerate(kf.split(data_features))
...     for fold_num, index_tup in enum_obj:
...         zip_obj = zip(["train", "test"], index_tup)
...         split_data = {}
...         for set_name, set_indices in zip_obj:
...             split_data[set_name] = (data_features, data_labels)
...         #x = {data_name:X.shape for data_name, (X,y) in split_data.items()}
...         #print(f"{data_name}: ", x)
...
...

```

```

...     for algorithm, accuracy in test_accuracy.items():
...         print(f"{algorithm} Test Accuracy: {accuracy * 100}")
...         accuracy_df = pd.DataFrame({
...             "data_set": [data_name],
...             "fold_id": [fold_num],
...             "algorithm": [algorithm],
...             "accuracy": [test_accuracy[algorithm]]})
...         accuracy_data_frames.append(accuracy_df)
...     print(f"*****End of
{data_name}{fold_num}*****")

```

zip: 2007 256

Best N-Neighbors = 1

Most Frequent Label = 0

Best Epoch: 6

Best Epoch: 7

KNeighborsClassifier + GridSearchCV Test Accuracy: 100.0

LogisticRegressionCV Test Accuracy: 97.50871948181366

TorchLearnerCV Linear Test Accuracy: 93.47284504235176

TorchLearnerCV Deep Test Accuracy: 97.45889387144993

featureless Test Accuracy: 17.887394120577977

\*\*\*\*\*End of zip(0)\*\*\*\*\*

zip: 2007 256

Best N-Neighbors = 1

Most Frequent Label = 0

Best Epoch: 7

Best Epoch: 6

KNeighborsClassifier + GridSearchCV Test Accuracy: 100.0

LogisticRegressionCV Test Accuracy: 97.50871948181366

TorchLearnerCV Linear Test Accuracy: 95.41604384653712

TorchLearnerCV Deep Test Accuracy: 96.81116093672148

featureless Test Accuracy: 17.887394120577977

\*\*\*\*\*End of zip(1)\*\*\*\*\*

zip: 2007 256

Best N-Neighbors = 1

Most Frequent Label = 0

Best Epoch: 7

Best Epoch: 9

KNeighborsClassifier + GridSearchCV Test Accuracy: 100.0

LogisticRegressionCV Test Accuracy: 97.50871948181366

TorchLearnerCV Linear Test Accuracy: 94.22022919780767

TorchLearnerCV Deep Test Accuracy: 97.85749875435974

featureless Test Accuracy: 17.887394120577977

\*\*\*\*\*End of zip(2)\*\*\*\*\*

MNIST: 10000 784

Best N-Neighbors = 3

Most Frequent Label = 1



Best Epoch: 18  
 Best Epoch: 11  
 KNeighborsClassifier + GridSearchCV Test Accuracy: 97.72999999999999  
 LogisticRegressionCV Test Accuracy: 94.57  
 TorchLearnerCV Linear Test Accuracy: 94.55  
 TorchLearnerCV Deep Test Accuracy: 99.99  
 featureless Test Accuracy: 11.35  
 \*\*\*\*\*End of MNIST(0)\*\*\*\*\*

MNIST: 10000 784  
 Best N-Neighbors = 3  
 Most Frequent Label = 1  
 Best Epoch: 17  
 Best Epoch: 8  
 KNeighborsClassifier + GridSearchCV Test Accuracy: 97.72999999999999  
 LogisticRegressionCV Test Accuracy: 94.57  
 TorchLearnerCV Linear Test Accuracy: 94.62  
 TorchLearnerCV Deep Test Accuracy: 99.7  
 featureless Test Accuracy: 11.35  
 \*\*\*\*\*End of MNIST(1)\*\*\*\*\*

MNIST: 10000 784  
 Best N-Neighbors = 3  
 Most Frequent Label = 1  
 Best Epoch: 14  
 Best Epoch: 8  
 KNeighborsClassifier + GridSearchCV Test Accuracy: 97.72999999999999  
 LogisticRegressionCV Test Accuracy: 94.57  
 TorchLearnerCV Linear Test Accuracy: 94.35  
 TorchLearnerCV Deep Test Accuracy: 99.47  
 featureless Test Accuracy: 11.35  
 \*\*\*\*\*End of MNIST(2)\*\*\*\*\*

```
>>> total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
```

```
>>> print(total_accuracy_df)
```


	data_set	fold_id	algorithm	accuracy
0	zip	0	KNeighborsClassifier + GridSearchCV	1.000000
1	zip	0	LogisticRegressionCV	0.975087
2	zip	0	TorchLearnerCV Linear	0.934728
3	zip	0	TorchLearnerCV Deep	0.974589
4	zip	0	featureless	0.178874
5	zip	1	KNeighborsClassifier + GridSearchCV	1.000000
6	zip	1	LogisticRegressionCV	0.975087
7	zip	1	TorchLearnerCV Linear	0.954160
8	zip	1	TorchLearnerCV Deep	0.968112
9	zip	1	featureless	0.178874
10	zip	2	KNeighborsClassifier + GridSearchCV	1.000000
11	zip	2	LogisticRegressionCV	0.975087

12	zip	2	TorchLearnerCV Linear	0.942202
13	zip	2	TorchLearnerCV Deep	0.978575
14	zip	2	featureless	0.178874
15	MNIST	0	KNeighborsClassifier + GridSearchCV	0.977300
16	MNIST	0	LogisticRegressionCV	0.945700
17	MNIST	0	TorchLearnerCV Linear	0.945500
18	MNIST	0	TorchLearnerCV Deep	0.999900
19	MNIST	0	featureless	0.113500
20	MNIST	1	KNeighborsClassifier + GridSearchCV	0.977300
21	MNIST	1	LogisticRegressionCV	0.945700
22	MNIST	1	TorchLearnerCV Linear	0.946200
23	MNIST	1	TorchLearnerCV Deep	0.997000
24	MNIST	1	featureless	0.113500
25	MNIST	2	KNeighborsClassifier + GridSearchCV	0.977300
26	MNIST	2	LogisticRegressionCV	0.945700
27	MNIST	2	TorchLearnerCV Linear	0.943500
28	MNIST	2	TorchLearnerCV Deep	0.994700
29	MNIST	2	featureless	0.113500

#### Test Loss Square Graph:

```
>>> gg = p9.ggplot(total_accuracy_df, p9.aes(x='accuracy', y='algorithm'))+\
...     p9.facet_grid('~data_set') + p9.geom_point()

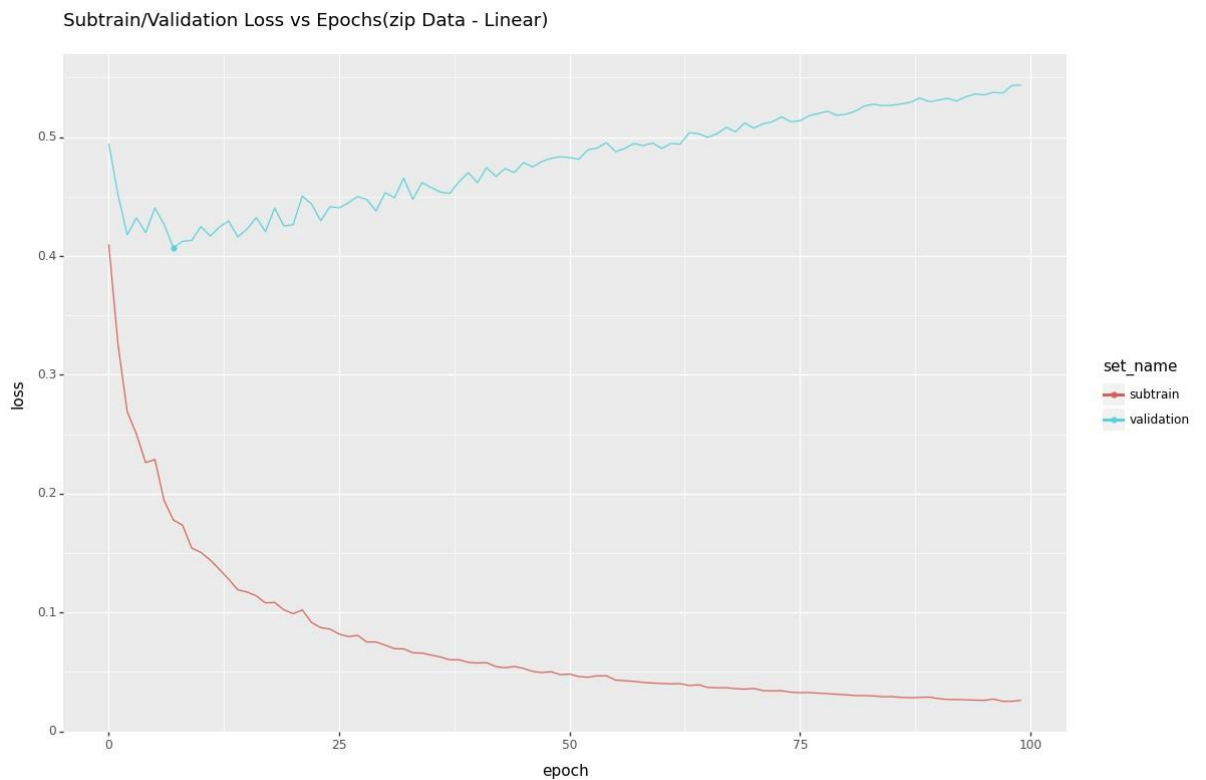
>>> gg.save("Test_square_loss.png", height = 8, width = 12)
```



Linear subtrain/validation loss graph (zip):

```
>>> gg1 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_loss["TorchLearnerCV Linear"])\n... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_min["min_df linear"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(zip Data - Linear)")
```


```
>>> gg1.save("Torch_validation_graph1.png", height = 8, width = 12)
```



Linear subtrain/validation loss graph (MNIST):

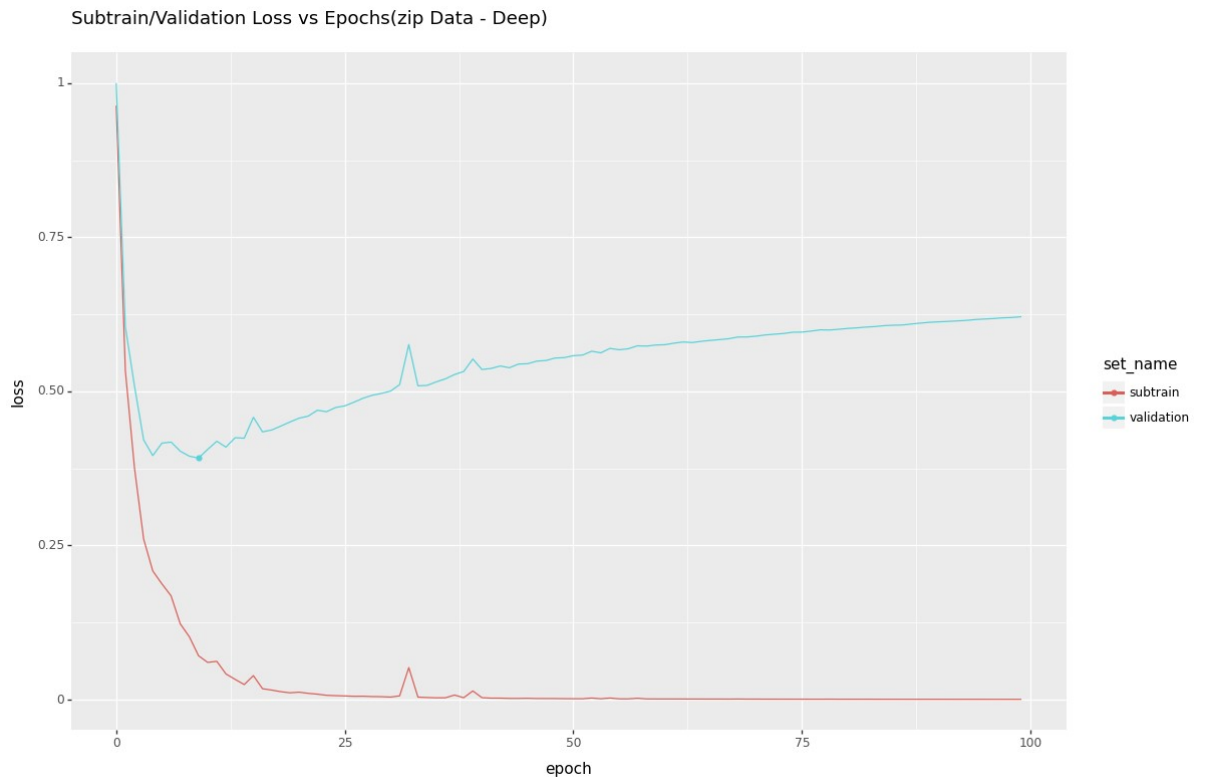
```
>>> gg3 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data = mnist_loss["TorchLearnerCV Linear"])\n... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = mnist_min["min_df linear"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(MNIST Data - Linear)")
```

```
>>> gg3.save("Torch_validation_graph3.png", height = 8, width = 12)
```



Deep subtrain/validation loss graph (zip):

```
>>> gg2 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_loss["TorchLearnerCV Deep"])\n... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = zip_min["min_df deep"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(zip Data - Deep)")\n\n>>> gg2.save("Torch_validation_graph2.png", height = 8, width = 12)
```



**Deep subtrain/validation loss graph (MNIST):**

```
>>> gg4 = p9.ggplot() + p9.geom_line(p9.aes(x='epoch', y='loss', color='set_name'), data = mnist_loss["TorchLearnerCV Deep"])\
... + p9.geom_point(p9.aes(x='epoch', y='loss', color='set_name'), data = mnist_min["min_df deep"]) + p9.ggtitle("Subtrain/Validation Loss vs Epochs(MNIST Data - Deep)")
```

```
>>> gg4.save("Torch_validation_graph4.png", height = 8, width = 12)
```

Subtrain/Validation Loss vs Epochs(MNIST Data - Deep)

