

CS599 (Deep Learning)

Homework – 13

1. Python Code:

```
import torch
import pandas as pd
import matplotlib
#matplotlib.use("agg")
import numpy as np
import plotnine as p9
import math
import pdb
```

```
from sklearn.model_selection import KFold, GridSearchCV, ParameterGrid
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from collections import Counter
```

```
data_set_dict = {"zip": ("zip.test.gz", 0)}
data_dict = {}
```

```
for data_name, (file_name, label_col_num) in data_set_dict.items():
    data_df = pd.read_csv(file_name, sep = " ", header = None)
    data_label_vec = data_df.iloc[:, label_col_num]
    is_01 = data_label_vec.isin([0, 1])
    data_01_df = data_df.loc[is_01, :]
    is_label_col = data_df.columns == label_col_num
    data_features = data_01_df.iloc[:, ~is_label_col]
    data_labels = data_01_df.iloc[:, is_label_col]
    data_dict[data_name] = (data_features, data_labels)
```

```
zip_df = pd.read_csv("zip.test.gz", sep = " ", header = None)
zip_label_col_num = 0
zip_label_vec = zip_df.iloc[:, zip_label_col_num]
is_71 = zip_label_vec.isin([7,1])
zip_71_df = zip_df.loc[is_71, :]
is_label_col = zip_71_df.columns == zip_label_col_num
zip_features = zip_71_df.iloc[:, ~is_label_col]
zip_labels = zip_71_df.iloc[:, is_label_col]
zip_labels = zip_labels.replace(7, 0)
data_dict["zip_71"] = (zip_features, zip_labels)
```

```
{data_name:X.shape for data_name, (X,y) in data_dict.items()}
```

```
{'zip': (623, 256), 'zip_71': (411, 256)}
```

```
class TorchModel(torch.nn.Module):
    def __init__(self, units_per_layer):
        super(TorchModel, self).__init__()
        self.conv = torch.nn.Sequential(
            torch.nn.Conv2d(1, 32, kernel_size=(3, 3)),
            torch.nn.Conv2d(32, 64, kernel_size=(3, 3)),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size = (3, 3)),
            torch.nn.Flatten(start_dim = 1))
        self.lin_seq = torch.nn.Sequential(
            torch.nn.Linear(1, 128),
            torch.nn.ReLU(),
            torch.nn.Linear(128, 2))
        self.fc = torch.nn.Sequential(
            self.conv,
            self.lin_seq)
```

```

def forward(self, features):
    _, flattened_size = self.conv(features[:1]).shape
    self.lin_seq[0] = torch.nn.Linear(flattened_size, 128)
    x = self.fc(features)
    return x

```

```

class TorchModel(torch.nn.Module):
    def __init__(self, units_per_layer):
        super(TorchModel, self).__init__()
        self.conv = torch.nn.Sequential(
            torch.nn.Conv2d(1, 32, kernel_size=(3, 3)),
            torch.nn.Conv2d(32, 64, kernel_size=(3, 3)),
            torch.nn.ReLU(),
            torch.nn.Flatten(start_dim = 1))
        self.lin_seq = torch.nn.Sequential(
            torch.nn.Linear(1, 128),
            torch.nn.ReLU(),
            torch.nn.Linear(128, 2))
        self.fc = torch.nn.Sequential(
            self.conv,
            self.lin_seq)

    def forward(self, features):
        _, flattened_size = self.conv(features[:1]).shape
        self.lin_seq[0] = torch.nn.Linear(flattened_size, 128)
        x = self.fc(features)
        return x

```

```

class CSV(torch.utils.data.Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels
    def __getitem__(self, item):
        return self.features[item,:], self.labels[item]
    def __len__(self):
        return len(self.labels)

```

```

class TorchLearner:
    def __init__(
        self, units_per_layer, opt_name, opt_params,
        batch_size=20, max_epochs=100):
        self.max_epochs = max_epochs
        self.batch_size=batch_size
        self.model = TorchModel(units_per_layer)
        self.loss_fun = torch.nn.CrossEntropyLoss()
        self.initial_step_size = 0.1
        self.end_step_size = 0.001
        self.last_step_number = 50
        self.opt_name = opt_name
        self.opt_params = opt_params
    def get_step_size(self, iteration):
        if iteration > self.last_step_number:
            return self.end_step_size
        prop_to_last_step = iteration/self.last_step_number
        return (1 - prop_to_last_step) * self.initial_step_size + \
            prop_to_last_step * self.end_step_size
    def fit(self, split_data_dict):
        ds = CSV(
            split_data_dict["subtrain"]["X"],
            split_data_dict["subtrain"]["y"])
        dl = torch.utils.data.DataLoader(
            ds, batch_size=self.batch_size, shuffle=True)
        train_df_list = []

        for epoch_number in range(self.max_epochs):
            step_size = self.get_step_size(epoch_number)
            if self.opt_name == "SGD":
                self.optimizer = torch.optim.SGD(self.model.parameters(), **self.opt_params, lr = step_size)
            elif self.opt_name == "Adam":
                self.optimizer = torch.optim.Adam(self.model.parameters(), **self.opt_params, lr = step_size)
            for batch_features, batch_labels in dl:
                nrow, ncol = batch_features.shape
                self.optimizer.zero_grad()
                loss_value = self.loss_fun(
                    self.model(batch_features.view(nrow, 1, 16, 16)), batch_labels)
                loss_value.backward()
                self.optimizer.step()

            for set_name, set_data in split_data_dict.items():
                n_row, n_col = set_data["X"].shape
                pred_vec = self.model(set_data["X"].reshape(n_row, 1, 16, 16))
                set_loss_value = self.loss_fun(pred_vec, set_data["y"])
                train_df_list.append(pd.DataFrame({
                    "set_name": [set_name],
                    "loss": float(set_loss_value),
                    "epoch": [epoch_number]
                }))
            self.train_df = pd.concat(train_df_list)
    def decision_function(self, test_features):
        with torch.no_grad():
            row, col = test_features.shape
            pred_vec = self.model(test_features.reshape(row, 1, 16, 16))
            return pred_vec

    def predict(self, test_features):
        pred_scores = self.decision_function(test_features)
        _, predicted = torch.max(pred_scores, 1)
        return predicted

```

```

class TorchLearnerCV:
    def __init__(self, n_folds, units_per_layer, opt_name = 'SGD', opt_params = {'momentum': 0.5}):
        self.units_per_layer = units_per_layer
        self.opt_name = opt_name
        self.opt_params = opt_params
        self.n_folds = n_folds

    def fit(self, train_features, train_labels):
        train_nrow, train_ncol = train_features.shape
        times_to_repeat = int(math.ceil(train_nrow/self.n_folds))
        fold_id_vec = np.tile(torch.arange(self.n_folds), times_to_repeat)[:train_nrow]
        np.random.shuffle(fold_id_vec)
        cv_data_list = []
        for validation_fold in range(self.n_folds):
            is_split = {
                "subtrain": fold_id_vec != validation_fold,
                "validation": fold_id_vec == validation_fold
            }
            split_data_dict = {}
            for set_name, is_set in is_split.items():
                set_y = train_labels[is_set]
                split_data_dict[set_name] = {
                    "X": train_features[is_set,:],
                    "y": set_y
                }
            learner = TorchLearner(self.units_per_layer, self.opt_name, self.opt_params)
            learner.fit(split_data_dict)
            cv_data_list.append(learner.train_df)
        self.cv_data = pd.concat(cv_data_list)
        self.train_df = self.cv_data.groupby(["set_name", "epoch"]).mean().reset_index()
        #print(self.train_df)
        valid_df = self.train_df.query("set_name=='validation'")
        #print(valid_df)
        best_epochs = valid_df["loss"].argmin()
        self.min_df = valid_df.query("epoch==%s"%(best_epochs))
        print("Best Epoch: ", best_epochs)
        #pdb.set_trace()
        self.final_learner = TorchLearner(self.units_per_layer, self.opt_name, self.opt_params, max_epochs=(best_epochs + 1))
        self.final_learner.fit({"subtrain": {"X": train_features, "y": train_labels}})
        return self.cv_data

    def predict(self, test_features):
        return self.final_learner.predict(test_features)

```

```

class MyCV:
    def __init__(self, estimator, param_grid, cv):
        """estimator: learner instance
        param_grid: list of dictionaries
        cv: number of folds"""
        self.cv = cv
        self.param_grid = param_grid
        self.estimator = estimator

    def fit_one(self, param_dict, X, y):
        """Run self.estimator.fit on one parameter combination"""
        for param_name, param_value in param_dict.items():
            #print(f"param_name = {param_name}, param_value = {param_value}")
            setattr(self.estimator, param_name, param_value)
        self.estimator.fit(X, y)

    def fit(self, X, y):
        """cross-validation for selecting the best dictionary is param_grid"""
        validation_df_list = []
        train_nrow, train_ncol = X.shape
        times_to_repeat = int(math.ceil(train_nrow/self.cv))
        fold_id_vec = np.tile(np.arange(self.cv), times_to_repeat)[:train_nrow]
        np.random.shuffle(fold_id_vec)
        for validation_fold in range(self.cv):
            is_split = {
                "subtrain": fold_id_vec != validation_fold,
                "validation": fold_id_vec == validation_fold
            }
            split_data_dict = {}
            for set_name, is_set in is_split.items():
                split_data_dict[set_name] = (
                    X[is_set],
                    y[is_set])
            for param_number, param_dict in enumerate(self.param_grid):
                self.fit_one(param_dict, *split_data_dict["subtrain"])
                X_valid, y_valid = split_data_dict["validation"]
                pred_valid = self.estimator.predict(X_valid)
                #pdb.set_trace()
                is_correct = pred_valid == y_valid
                #self.estimator.fit(*split_data_dict["validation"])
                valid_loss = self.estimator.train_df.query("set_name=='validation'")["loss"].mean()
                subtrain_loss = self.estimator.train_df.query("set_name=='subtrain'")["loss"].mean()
                validation_row1 = pd.DataFrame({
                    "set_name": "subtrain",
                    "validation_fold": validation_fold,
                    "accuracy_percent": float(is_correct.float().mean()),
                    "param_number": [param_number],
                    "loss": float(subtrain_loss)
                }, index = [0])
                validation_row2 = pd.DataFrame({
                    "set_name": "validation",
                    "validation_fold": validation_fold,
                    "accuracy_percent": float(is_correct.float().mean()),
                    "param_number": [param_number],
                    "loss": float(valid_loss)
                }, index = [0])
                validation_df_list.append(validation_row1)
                validation_df_list.append(validation_row2)
            self.validation_df = pd.concat(validation_df_list)
            self.mean_valid_loss = self.validation_df.groupby("param_number")["loss"].mean().reset_index()
            self.train_df = self.validation_df.groupby(["set_name", "loss"]).mean().reset_index()
            best_index = self.mean_valid_loss["loss"].argmin()
            #pdb.set_trace()
            valid_df = self.train_df.query("set_name == 'validation'")
            self.min_df = valid_df.query("param_number==%s"%(best_index))
            self.best_param_dict = self.param_grid[best_index]
            self.fit_one(self.best_param_dict, X, y)

    def predict(self, X):
        return self.estimator.predict(X)

```

```

accuracy_data_frames = []
loss_data_dict = {}
min_df_dict = {}
for data_name, (data_features, data_labels) in data_dict.items():
    kf = KFold(n_splits=3, shuffle=True, random_state=3)
    enum_obj = enumerate(kf.split(data_features))
    for fold_num, index_tup in enum_obj:
        zip_obj = zip(["train", "test"], index_tup)
        split_data = {}
        for set_name, set_indices in zip_obj:
            split_data[set_name] = (torch.from_numpy(data_features.iloc[set_indices, :].to_numpy()).float(),
                                    torch.from_numpy(np.ravel(data_labels.iloc[set_indices])).flatten())
        x = {data_name:X.shape for data_name, (X,y) in split_data.items()}
        print(f"{data_name}: ", x)
        train_features, train_labels = split_data["train"]
        nrow, ncol = train_features.shape
        print(f"{data_name}: ", nrow, ncol)
        test_features, test_labels = split_data["test"]

        #kneighbors
        knn = KNeighborsClassifier()
        hp_parameters = {"n_neighbors": list(range(1, 21))}
        grid = GridSearchCV(knn, hp_parameters, cv=3)
        grid.fit(train_features, train_labels)
        best_n_neighbors = grid.best_params_['n_neighbors']
        print("Best N-Neighbors = ", best_n_neighbors)
        knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
        knn.fit(train_features, train_labels)
        knn_pred = knn.predict(test_features)

        #linear model
        pipe = make_pipeline(StandardScaler(), LogisticRegressionCV(cv=3, max_iter=2000))
        pipe.fit(train_features, train_labels)
        lr_pred = pipe.predict(test_features)

        #Featureless
        y_train_series = pd.Series(train_labels)
        # create a featureless baseline
        most_frequent_label = y_train_series.value_counts().idxmax()
        print("Most Frequent Label = ", most_frequent_label)
        featureless_pred = np.repeat(most_frequent_label, len(test_features))

    param_grid = []
    for momentum in 0.1, 0.5:
        param_grid.append({
            "opt_name": "SGD",
            "opt_params": {"momentum": momentum}
        })

```

```

    })
    for beta1 in 0.85, 0.9, 0.95:
        for beta2 in 0.99, 0.999, 0.9999:
            param_grid.append({
                "opt_name": "Adam",
                "opt_params": {"betas": (beta1, beta2)}
            })

#TorchLearnerCV + Deep
conv_learner = MyCV(
    estimator = TorchLearnerCV(3, [ncol, 64, 32, 128, 1]),
    param_grid = param_grid,
    cv = 2)
conv_learner.fit(train_features, train_labels)
print(f"Best param_dict: {conv_learner.best_param_dict}")
cl_pred = conv_learner.predict(test_features)

min_df_dict[data_name] = {'min_df': conv_learner.estimator.min_df}

loss_data_dict[data_name] = {'conv_learner': conv_learner.estimator.train_df}

# store predict data in dict
pred_dict = {'KNeighborsClassifier + GridSearchCV': knn_pred,
             'LogisticRegressionCV': lr_pred,
             'ConvolutionalMLP': cl_pred,
             'featureless': featureless_pred}

test_accuracy = {}
for algorithm, predictions in pred_dict.items():
    accuracy = accuracy_score(test_labels, predictions)
    test_accuracy[algorithm] = accuracy

for algorithm, accuracy in test_accuracy.items():
    print(f"{algorithm} Test Accuracy: {accuracy * 100}")
    accuracy_df = pd.DataFrame({
        "data_set": [data_name],
        "fold_id": [fold_num],
        "algorithm": [algorithm],
        "accuracy": [test_accuracy[algorithm]]})
    accuracy_data_frames.append(accuracy_df)
    print(f"*****End of {data_name}({fold_num})*****")

```

```

total_accuracy_df = pd.concat(accuracy_data_frames, ignore_index = True)
print(total_accuracy_df)

```

```

gg = p9.ggplot(total_accuracy_df, p9.aes(x = 'accuracy', y = 'algorithm'))+\
    p9.facet_grid('~data_set') + p9.geom_point()

```

```

print(gg)

```

```
gg1 = p9.ggplot() +\
  p9.geom_line(\
    p9.aes(\
      x = "epoch",\
      y = "loss",\
      color = "set_name"\
    ),\
    data = loss_data_dict["zip"]['conv_learner']) +\
  p9.geom_point(\
    p9.aes(\
      x = "epoch",\
      y = "loss",\
      color = "set_name"\
    ),\
    data = min_df_dict["zip"]['min_df']) +\
  p9.ggtitle("Subtrain/Validation Loss vs Epochs(Zip - Data) with Max Pooling")
```

```
print(gg1)
```

```
gg2 = p9.ggplot() +\
  p9.geom_line(\
    p9.aes(\
      x = "epoch",\
      y = "loss",\
      color = "set_name"\
    ),\
    data = loss_data_dict["zip_71"]['conv_learner']) +\
  p9.geom_point(\
    p9.aes(\
      x = "epoch",\
      y = "loss",\
      color = "set_name"\
    ),\
    data = min_df_dict["zip_71"]['min_df']) +\
  p9.ggtitle("Subtrain/Validation Loss vs Epochs(zip_71 - Data) with Max Pooling")
```

```
print(gg2)
```


2. Output:

Accuracy with Max Pooling:

```
zip: {'train': torch.Size([415, 256]), 'test': torch.Size([208, 256])}
zip: 415 256
Best N-Neighbors = 1
Most Frequent Label = 0
Best Epoch: 88
Best Epoch: 75
Best Epoch: 87
Best Epoch: 43
Best Epoch: 92
Best Epoch: 79
Best Epoch: 89
Best Epoch: 96
Best Epoch: 31
Best Epoch: 31
Best Epoch: 27
Best Epoch: 82
Best Epoch: 43
Best Epoch: 75
Best Epoch: 37
Best Epoch: 75
Best Epoch: 63
Best Epoch: 97
Best Epoch: 47
Best Epoch: 90
Best Epoch: 86
Best Epoch: 22
Best Epoch: 43
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 100.0
LogisticRegressionCV Test Accuracy: 99.51923076923077
ConvolutionalMLP Test Accuracy: 96.63461538461539
featureless Test Accuracy: 58.65384615384615
*****End of zip(0)*****
zip: {'train': torch.Size([415, 256]), 'test': torch.Size([208, 256])}
zip: 415 256
Best N-Neighbors = 1
Most Frequent Label = 0
Best Epoch: 80
Best Epoch: 72
Best Epoch: 21
Best Epoch: 73
Best Epoch: 15
Best Epoch: 66
Best Epoch: 87
Best Epoch: 24
Best Epoch: 35
Best Epoch: 94
Best Epoch: 32
Best Epoch: 49
Best Epoch: 77
Best Epoch: 50
Best Epoch: 54
Best Epoch: 29
Best Epoch: 44
Best Epoch: 60
Best Epoch: 55
Best Epoch: 29
Best Epoch: 47
Best Epoch: 31
Best Epoch: 26
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.51923076923077
LogisticRegressionCV Test Accuracy: 99.03846153846155
ConvolutionalMLP Test Accuracy: 98.07692307692307
featureless Test Accuracy: 57.21153846153846
*****End of zip(1)*****
```

```

zip: {'train': torch.Size([416, 256]), 'test': torch.Size([207, 256])}
zip: 416 256
Best N-Neighbors = 4
Most Frequent Label = 0
Best Epoch: 95
Best Epoch: 93
Best Epoch: 27
Best Epoch: 20
Best Epoch: 37
Best Epoch: 50
Best Epoch: 20
Best Epoch: 48
Best Epoch: 19
Best Epoch: 21
Best Epoch: 28
Best Epoch: 75
Best Epoch: 91
Best Epoch: 61
Best Epoch: 80
Best Epoch: 40
Best Epoch: 94
Best Epoch: 73
Best Epoch: 26
Best Epoch: 36
Best Epoch: 39
Best Epoch: 39
Best Epoch: 59
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.03381642512076
LogisticRegressionCV Test Accuracy: 99.03381642512076
ConvolutionalMLP Test Accuracy: 98.06763285024155
featureless Test Accuracy: 57.00483091787439
*****End of zip(2)*****
zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 1
Most Frequent Label = 1
Best Epoch: 61
Best Epoch: 89
Best Epoch: 99
Best Epoch: 89
Best Epoch: 41
Best Epoch: 34
Best Epoch: 52
Best Epoch: 38
Best Epoch: 46
Best Epoch: 98
Best Epoch: 91
Best Epoch: 63
Best Epoch: 68
Best Epoch: 63
Best Epoch: 84
Best Epoch: 54
Best Epoch: 65
Best Epoch: 67
Best Epoch: 63
Best Epoch: 48
Best Epoch: 39
Best Epoch: 83
Best Epoch: 71
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.27007299270073
LogisticRegressionCV Test Accuracy: 99.27007299270073
ConvolutionalMLP Test Accuracy: 96.35036496350365
featureless Test Accuracy: 62.77372262773723
*****End of zip_71(0)*****


```

```


zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 1
Most Frequent Label = 1
Best Epoch: 87
Best Epoch: 97
Best Epoch: 82
Best Epoch: 46
Best Epoch: 45
Best Epoch: 81
Best Epoch: 98
Best Epoch: 42
Best Epoch: 50
Best Epoch: 61
Best Epoch: 30
Best Epoch: 49
Best Epoch: 89
Best Epoch: 33
Best Epoch: 75
Best Epoch: 68
Best Epoch: 49
Best Epoch: 57
Best Epoch: 69
Best Epoch: 71
Best Epoch: 96
Best Epoch: 96
Best Epoch: 86
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 97.8102189781022
LogisticRegressionCV Test Accuracy: 96.35036496350365
ConvolutionalMLP Test Accuracy: 95.62043795620438
featureless Test Accuracy: 62.77372262773723
*****End of zip_71(1)*****
zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 2
Most Frequent Label = 1
Best Epoch: 95
Best Epoch: 97
Best Epoch: 91
Best Epoch: 32
Best Epoch: 84
Best Epoch: 56
Best Epoch: 44
Best Epoch: 33
Best Epoch: 89
Best Epoch: 62
Best Epoch: 80
Best Epoch: 43
Best Epoch: 94
Best Epoch: 83
Best Epoch: 48
Best Epoch: 44
Best Epoch: 75
Best Epoch: 90
Best Epoch: 95
Best Epoch: 79
Best Epoch: 45
Best Epoch: 91
Best Epoch: 86
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.27007299270073
LogisticRegressionCV Test Accuracy: 99.27007299270073
ConvolutionalMLP Test Accuracy: 94.8905109489051
featureless Test Accuracy: 67.15328467153284
*****End of zip_71(2)*****

```


	data_set	fold_id	algorithm	accuracy
0	zip	0	KNeighborsClassifier + GridSearchCV	1.000000
1	zip	0	LogisticRegressionCV	0.995192
2	zip	0	ConvolutionalMLP	0.966346
3	zip	0	featureless	0.586538
4	zip	1	KNeighborsClassifier + GridSearchCV	0.995192
5	zip	1	LogisticRegressionCV	0.990385
6	zip	1	ConvolutionalMLP	0.980769
7	zip	1	featureless	0.572115
8	zip	2	KNeighborsClassifier + GridSearchCV	0.990338
9	zip	2	LogisticRegressionCV	0.990338
10	zip	2	ConvolutionalMLP	0.980676
11	zip	2	featureless	0.570048
12	zip_71	0	KNeighborsClassifier + GridSearchCV	0.992701
13	zip_71	0	LogisticRegressionCV	0.992701
14	zip_71	0	ConvolutionalMLP	0.963504
15	zip_71	0	featureless	0.627737
16	zip_71	1	KNeighborsClassifier + GridSearchCV	0.978102
17	zip_71	1	LogisticRegressionCV	0.963504
18	zip_71	1	ConvolutionalMLP	0.956204
19	zip_71	1	featureless	0.627737
20	zip_71	2	KNeighborsClassifier + GridSearchCV	0.992701
21	zip_71	2	LogisticRegressionCV	0.992701
22	zip_71	2	ConvolutionalMLP	0.948905
23	zip_71	2	featureless	0.671533



Subtrain/Validation Loss vs Epochs(Zip - Data) with Max Pooling



Subtrain/Validation Loss vs Epochs(zip_71 - Data) with Max Pool



Accuracy without Max Pooling:

```
zip: {'train': torch.Size([415, 256]), 'test': torch.Size([208, 256])}
zip: 415 256
Best N-Neighbors = 1
Most Frequent Label = 0
Best Epoch: 63
Best Epoch: 50
Best Epoch: 31
Best Epoch: 53
Best Epoch: 98
Best Epoch: 48
Best Epoch: 27
Best Epoch: 60
Best Epoch: 37
Best Epoch: 96
Best Epoch: 35
Best Epoch: 52
Best Epoch: 61
Best Epoch: 51
Best Epoch: 82
Best Epoch: 25
Best Epoch: 31
Best Epoch: 95
Best Epoch: 68
Best Epoch: 94
Best Epoch: 27
Best Epoch: 66
Best Epoch: 91
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 100.0
LogisticRegressionCV Test Accuracy: 99.51923076923077
ConvolutionalMLP Test Accuracy: 100.0
featureless Test Accuracy: 58.65384615384615
*****End of zip(0)*****
zip: {'train': torch.Size([415, 256]), 'test': torch.Size([208, 256])}
zip: 415 256
Best N-Neighbors = 1
Most Frequent Label = 0
Best Epoch: 94
Best Epoch: 67
Best Epoch: 38
Best Epoch: 75
Best Epoch: 81
Best Epoch: 38
Best Epoch: 99
Best Epoch: 97
Best Epoch: 38
Best Epoch: 21
Best Epoch: 31
Best Epoch: 81
Best Epoch: 73
Best Epoch: 42
Best Epoch: 44
Best Epoch: 37
Best Epoch: 52
Best Epoch: 54
Best Epoch: 89
Best Epoch: 96
Best Epoch: 65
Best Epoch: 60
Best Epoch: 48
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.51923076923077
LogisticRegressionCV Test Accuracy: 99.03846153846155
ConvolutionalMLP Test Accuracy: 97.11538461538461
featureless Test Accuracy: 57.21153846153846
*****End of zip(1)*****
```

```

zip: {'train': torch.Size([416, 256]), 'test': torch.Size([207, 256])}
zip: 416 256
Best N-Neighbors = 4
Most Frequent Label = 0
Best Epoch: 75
Best Epoch: 69
Best Epoch: 58
Best Epoch: 34
Best Epoch: 27
Best Epoch: 42
Best Epoch: 21
Best Epoch: 79
Best Epoch: 59
Best Epoch: 73
Best Epoch: 31
Best Epoch: 63
Best Epoch: 58
Best Epoch: 79
Best Epoch: 68
Best Epoch: 28
Best Epoch: 39
Best Epoch: 31
Best Epoch: 29
Best Epoch: 54
Best Epoch: 92
Best Epoch: 18
Best Epoch: 86
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.03381642512076
LogisticRegressionCV Test Accuracy: 99.03381642512076
ConvolutionalMLP Test Accuracy: 97.58454106280193
featureless Test Accuracy: 57.00483091787439
*****End of zip(2)*****
zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 1
Most Frequent Label = 1
Best Epoch: 91
Best Epoch: 79
Best Epoch: 64
Best Epoch: 40
Best Epoch: 95
Best Epoch: 60
Best Epoch: 37
Best Epoch: 85
Best Epoch: 42
Best Epoch: 91
Best Epoch: 71
Best Epoch: 90
Best Epoch: 75
Best Epoch: 73
Best Epoch: 57
Best Epoch: 31
Best Epoch: 65
Best Epoch: 49
Best Epoch: 84
Best Epoch: 34
Best Epoch: 87
Best Epoch: 50
Best Epoch: 91
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.27007299270073
LogisticRegressionCV Test Accuracy: 99.27007299270073
ConvolutionalMLP Test Accuracy: 97.8102189781022
featureless Test Accuracy: 62.77372262773723
*****End of zip_71(0)*****

```


```

zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 1
Most Frequent Label = 1
Best Epoch: 90
Best Epoch: 94
Best Epoch: 43
Best Epoch: 96
Best Epoch: 60
Best Epoch: 23
Best Epoch: 41
Best Epoch: 85
Best Epoch: 75
Best Epoch: 86
Best Epoch: 56
Best Epoch: 87
Best Epoch: 82
Best Epoch: 93
Best Epoch: 93
Best Epoch: 66
Best Epoch: 95
Best Epoch: 94
Best Epoch: 51
Best Epoch: 88
Best Epoch: 98
Best Epoch: 82
Best Epoch: 98
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 97.8102189781022
LogisticRegressionCV Test Accuracy: 96.35036496350365
ConvolutionalMLP Test Accuracy: 94.16058394160584
featureless Test Accuracy: 62.77372262773723
*****End of zip_71(1)*****


zip_71: {'train': torch.Size([274, 256]), 'test': torch.Size([137, 256])}
zip_71: 274 256
Best N-Neighbors = 2
Most Frequent Label = 1
Best Epoch: 90
Best Epoch: 43
Best Epoch: 98
Best Epoch: 52
Best Epoch: 75
Best Epoch: 33
Best Epoch: 53
Best Epoch: 92
Best Epoch: 71
Best Epoch: 65
Best Epoch: 59
Best Epoch: 95
Best Epoch: 29
Best Epoch: 99
Best Epoch: 52
Best Epoch: 53
Best Epoch: 41
Best Epoch: 46
Best Epoch: 89
Best Epoch: 92
Best Epoch: 78
Best Epoch: 94
Best Epoch: 68
Best param_dict: {'opt_name': 'SGD', 'opt_params': {'momentum': 0.5}}
KNeighborsClassifier + GridSearchCV Test Accuracy: 99.27007299270073
LogisticRegressionCV Test Accuracy: 99.27007299270073
ConvolutionalMLP Test Accuracy: 97.8102189781022
featureless Test Accuracy: 67.15328467153284
*****End of zip_71(2)*****

```


	data_set	fold_id	algorithm	accuracy
0	zip	0	KNeighborsClassifier + GridSearchCV	1.000000
1	zip	0	LogisticRegressionCV	0.995192
2	zip	0	ConvolutionalMLP	1.000000
3	zip	0	featureless	0.586538
4	zip	1	KNeighborsClassifier + GridSearchCV	0.995192
5	zip	1	LogisticRegressionCV	0.990385
6	zip	1	ConvolutionalMLP	0.971154
7	zip	1	featureless	0.572115
8	zip	2	KNeighborsClassifier + GridSearchCV	0.990338
9	zip	2	LogisticRegressionCV	0.990338
10	zip	2	ConvolutionalMLP	0.975845
11	zip	2	featureless	0.570048
12	zip_71	0	KNeighborsClassifier + GridSearchCV	0.992701
13	zip_71	0	LogisticRegressionCV	0.992701
14	zip_71	0	ConvolutionalMLP	0.978102
15	zip_71	0	featureless	0.627737
16	zip_71	1	KNeighborsClassifier + GridSearchCV	0.978102
17	zip_71	1	LogisticRegressionCV	0.963504
18	zip_71	1	ConvolutionalMLP	0.941606
19	zip_71	1	featureless	0.627737
20	zip_71	2	KNeighborsClassifier + GridSearchCV	0.992701
21	zip_71	2	LogisticRegressionCV	0.992701
22	zip_71	2	ConvolutionalMLP	0.978102
23	zip_71	2	featureless	0.671533



Subtrain/Validation Loss vs Epochs(Zip - Data)



Subtrain/Validation Loss vs Epochs(zip_71 - Data)

